

Big data science

Day 4

F. Legger - INFN Torino

<https://github.com/Course-bigDataAndML/MLCourse-2425>

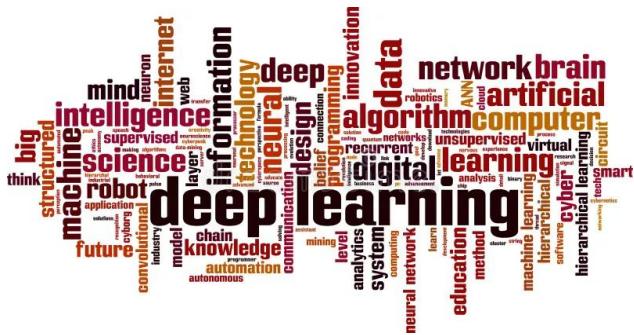
- So long, and thanks for all the images!
 - Taken freely from the web
 - Credits go to original creators



We Learned

- Big data, analytics
 - Distributed computing
 - ML
 - Supervised and unsupervised models
 - Evaluate performances

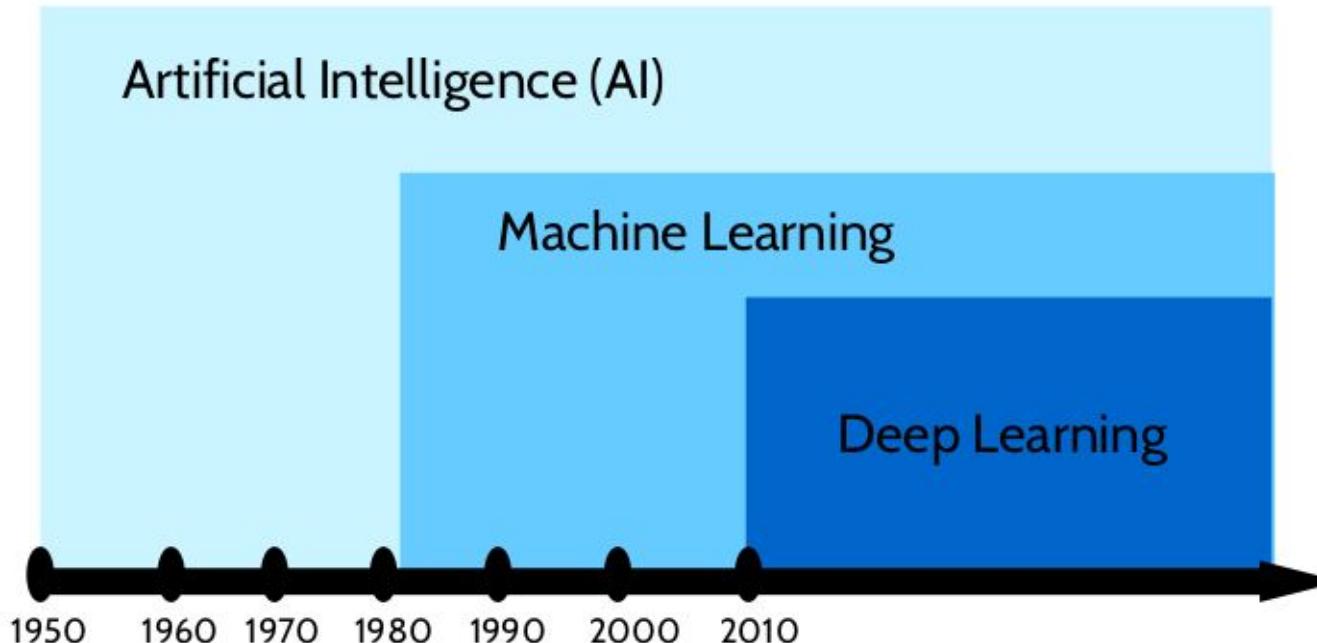
Today



- Deep learning
 - neural networks
 - brief history
 - main architectures
 - how to train

*Deep Learning is a subfield of ML concerned with algorithms inspired by the structure and function of the brain called **artificial neural networks***

[Jason Brownlee]

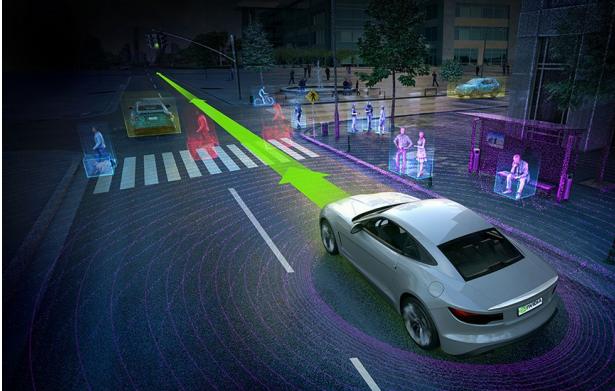
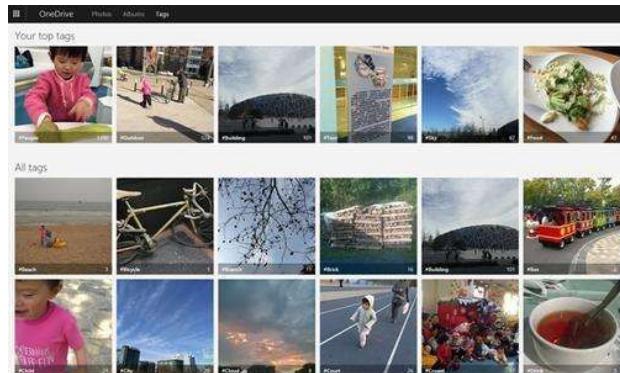


Machine translation

Real-time translation into Mandarin Chinese (2012)



Visual recognition, ResNet better than humans (2015)



Self driving cars, Tesla autopilot (2014)



Chatbots, ChatGPT (2022)

Creativity

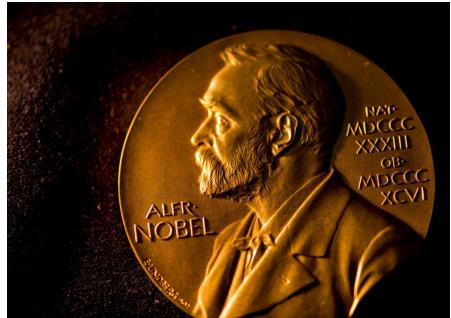
NST (2015), GAN (2014)



Strategy games DeepMind beats Go world champion (2017)



And the winner is...



- The Nobel Prize in **Physics 2024** was awarded jointly to **John J. Hopfield** (Princeton) and **Geoffrey E. Hinton** (U. Toronto) "for foundational discoveries and inventions that enable machine learning with artificial neural networks"
- The Nobel Prize in **Chemistry 2024** was divided, one half awarded to **David Baker** (U. Washington) "for computational protein design", the other half jointly to **Demis Hassabis** (Google DeepMind) and **John Jumper** (Google DeepMind) "for protein structure prediction"

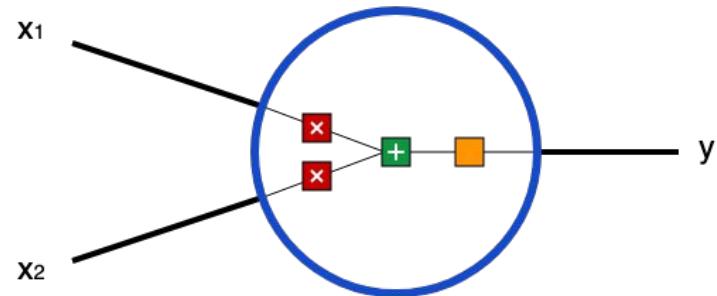
Short recap: Neuron

- each **input x** is multiplied by a **weight w**

$$x_1 \rightarrow x_1 * w_1$$



$$x_2 \rightarrow x_2 * w_2$$



- all the weighted inputs are added together with a **bias b**

$$(x_1 * w_1) + (x_2 * w_2) + b$$



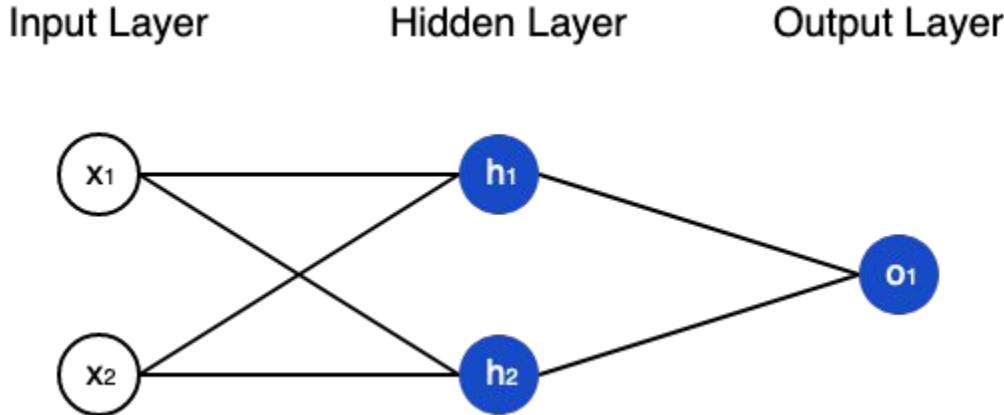
- the sum is passed through an **activation function f**



$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$

Neural network

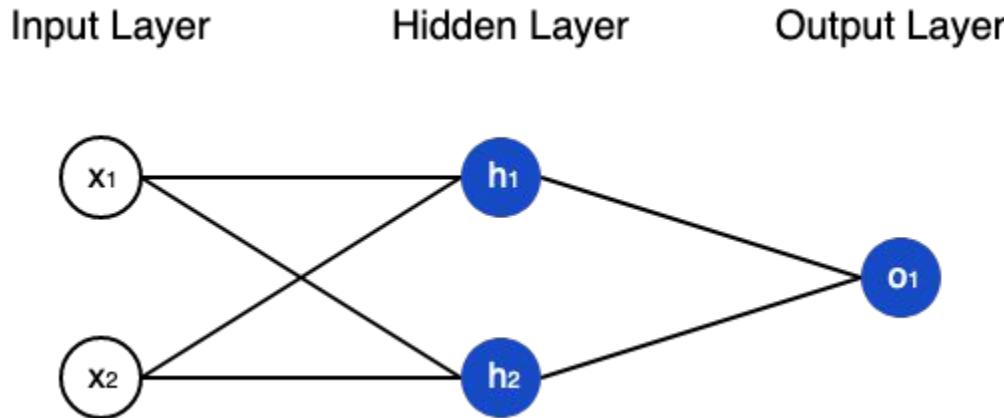
- built by combining neurons into layers
- a **hidden layer** is any layer between the **input** (first) layer and **output** (last) layer
 - there can be multiple hidden layers
- **Feedforward:** process of passing inputs forward to get an output



Neural network

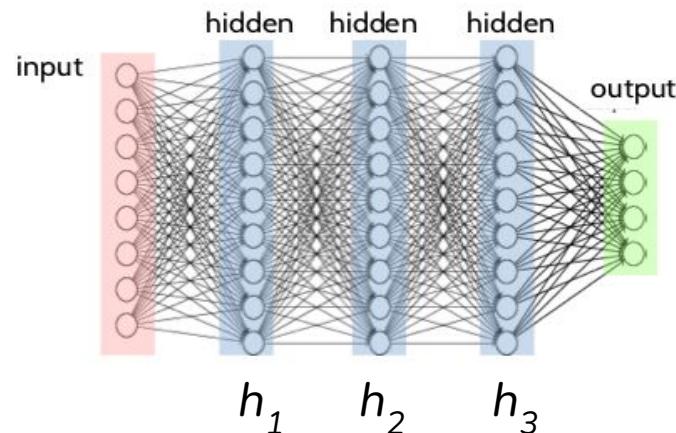
This network has:

- one **input** layer with 2 inputs
- one **hidden** layer with 2 neurons
- one **output** layer with 1 neuron



Deep Learning

- Neural network with several layers
 - Deep vs shallow
- A family of parametric models which learn non-linear hierarchical representations:

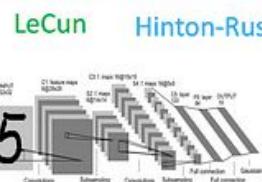
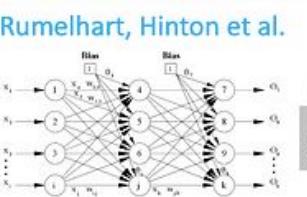
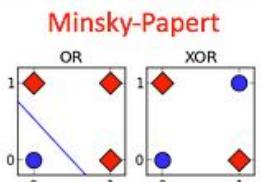
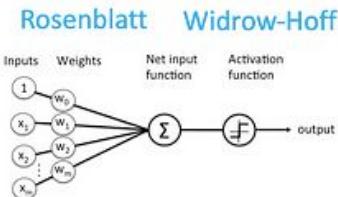
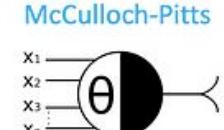
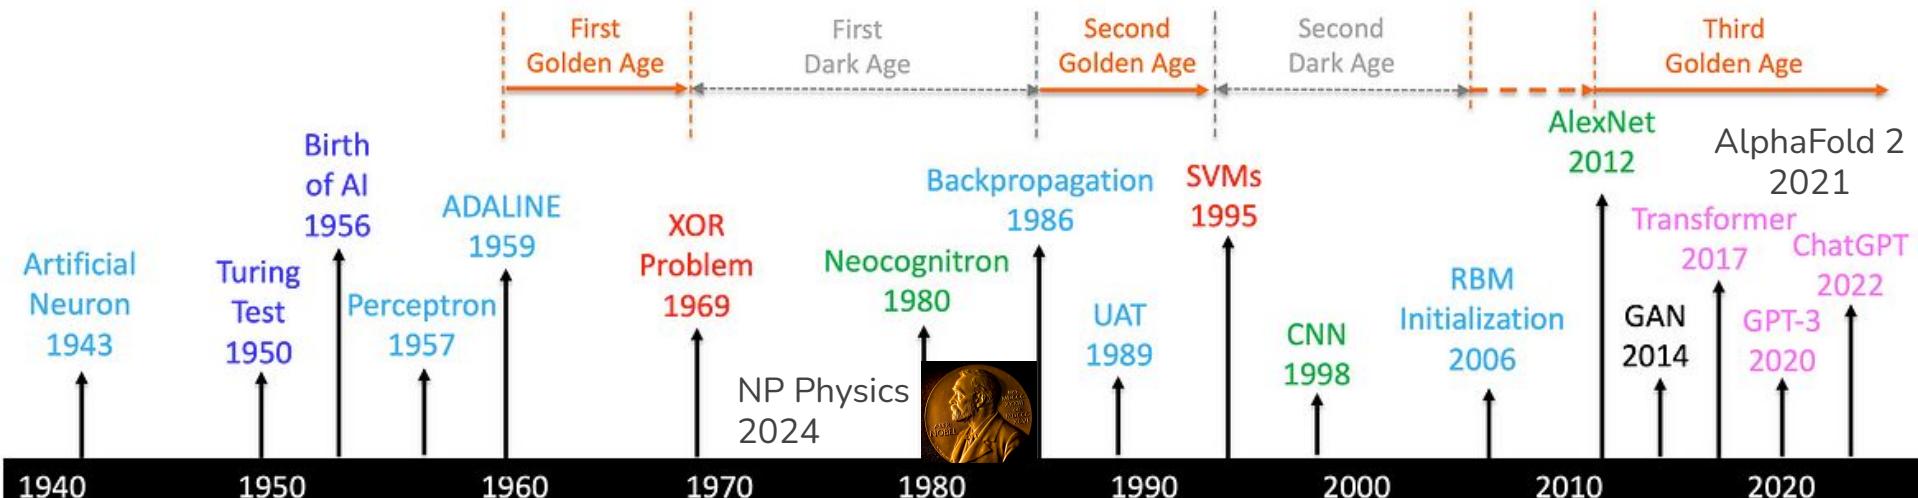


$$a_L(\mathbf{x}; \Theta) = h_L(h_{L-1}(\dots(h_1(\mathbf{x}, \theta_1), \theta_{L-1}), \theta_L)$$

input parameters of the network non-linear activation function parameters of layer L

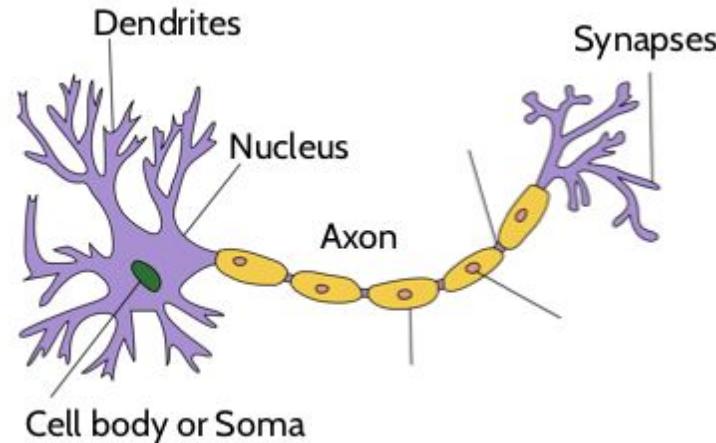
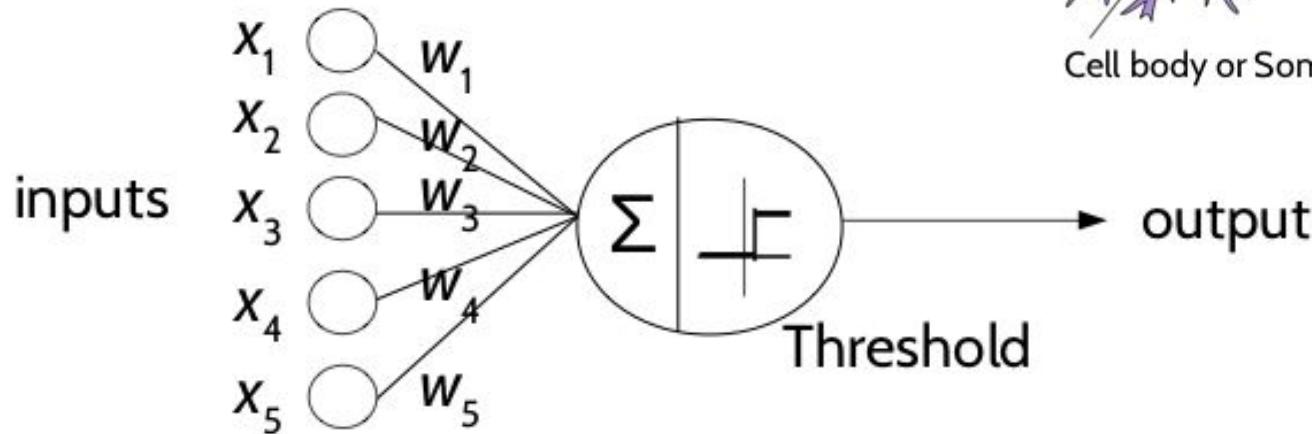
A Brief History of AI with Deep Learning

NP Chemistry
2024



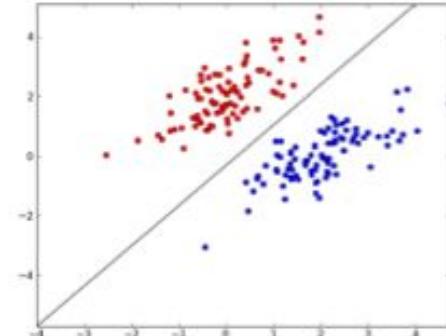
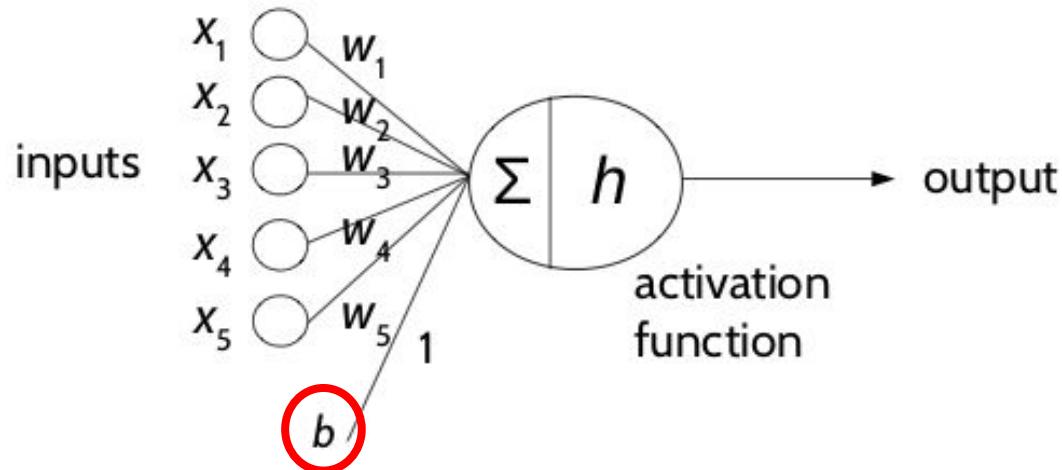
1943: McCulloch & Pitts Model

- Early model of artificial neuron
- Generates a binary output
- The weights values are fixed



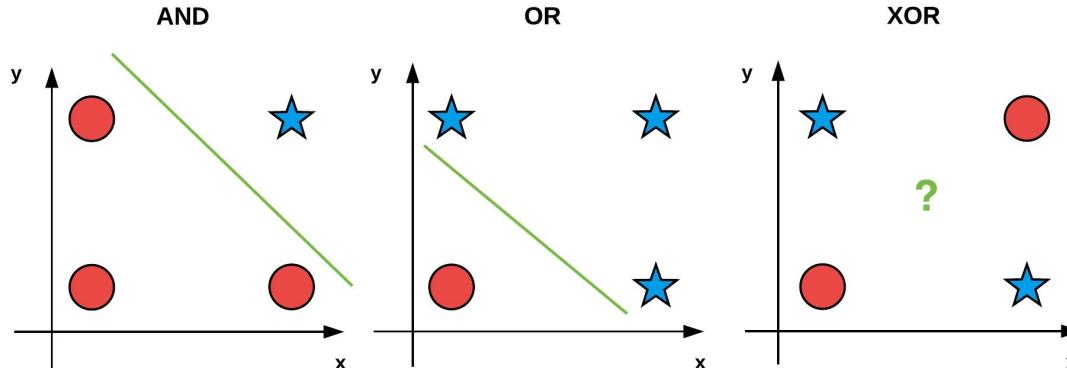
1958: Perceptron by Rosenblatt

- Perceptron as a machine for linear classification
- Main idea: Learn the weights and consider bias.
 - One weight per input
 - Multiply weights with respective inputs and add bias
 - If result larger than **threshold** return 1, otherwise 0



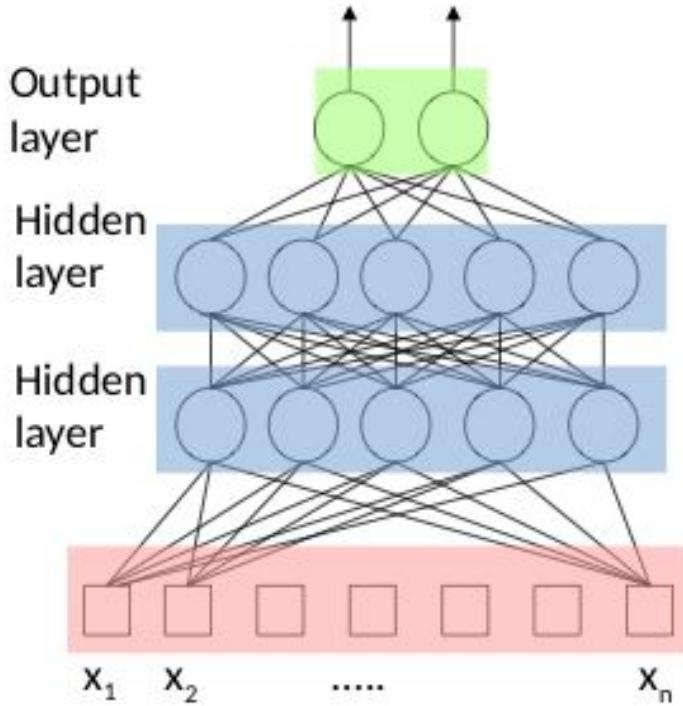
1970: First NN winter

- Minsky showed that the XOR cannot be solved by perceptrons
- Single-layer perceptrons are only capable of learning linearly separable patterns
- Neural models cannot be applied to complex tasks



1980: Multi-layer Neural Network

- 1980s. Multi-layer Perceptrons (MLP) can solve XOR.
- ML Feed Forward Neural Networks:
 - Densely connect artificial neurons to realize compositions of non-linear functions
 - The information is propagated from the inputs to the outputs
 - The input data are usually n -dimensional feature vectors
 - Tasks: Classification, Regression



How to train it?

- Rosenblatt algorithm* not applicable, as it expects to know the desired target
 - For hidden layers we cannot know the desired target
- Learning MLP for complicated functions can be solved with **Back propagation** (1986)**
 - efficient algorithm for complex NN which processes large training sets

* Remember? Rosenblatt developed a method to train a single neuron

** later today

1989: Universal Approximation Theorem (UAT)

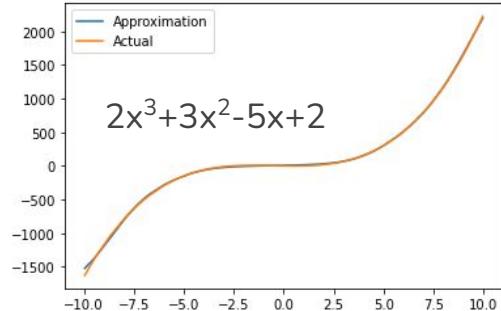
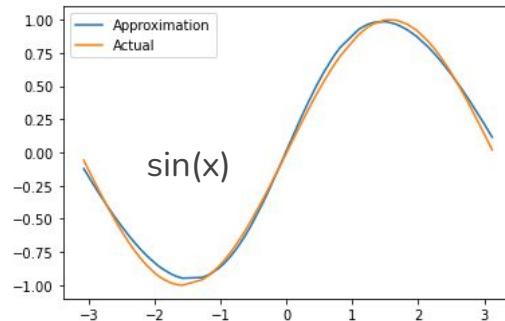
- A multilayered neural network with a **single hidden layer** can approximate any continuous function to any desired precision

Theorem 1. Let σ be any continuous discriminatory function. Then finite sums of the form

$$G(x) = \sum_{j=1}^N \alpha_j \sigma(y_j^T x + \theta_j) \quad (2)$$

are dense in $C(I_n)$. In other words, given any $f \in C(I_n)$ and $\varepsilon > 0$, there is a sum, $G(x)$, of the above form, for which

$$|G(x) - f(x)| < \varepsilon \quad \text{for all } x \in I_n.$$

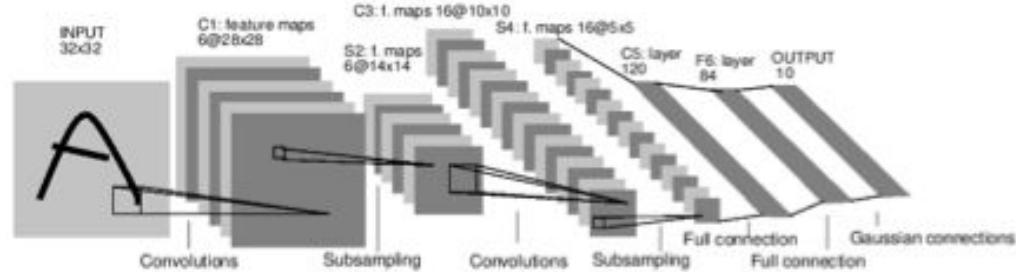
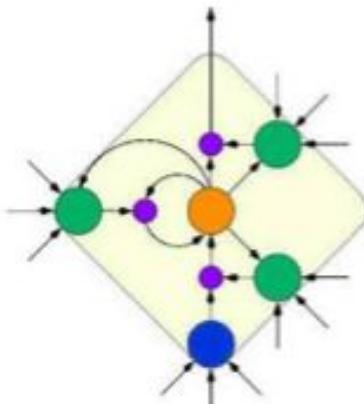


Cybenko, G. (1989). "Approximation by superpositions of a sigmoidal function". *Mathematics of Control, Signals, and Systems*. 2 (4): 303–314
Hornik, Kurt; Stinchcombe, Maxwell; White, Halbert (January 1989). "Multilayer feedforward networks are universal approximators". *Neural Networks*. 2 (5): 359–366

1990s: CNN and LSTM

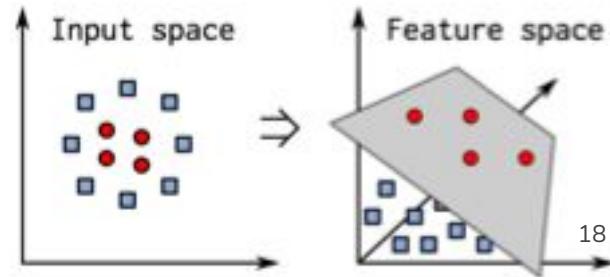
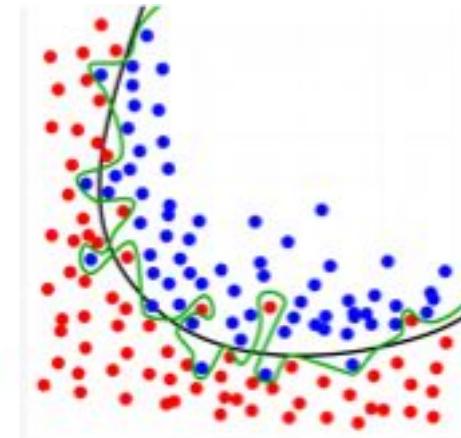
- Important advances in the field:
 - Backpropagation
 - Recurrent Long-Short Term Memory Networks (Schmidhuber, 1997)
 - Convolutional Neural Networks - LeNet: OCR solved before 2000s (LeCun, 1998).

OCR: Optical character recognition



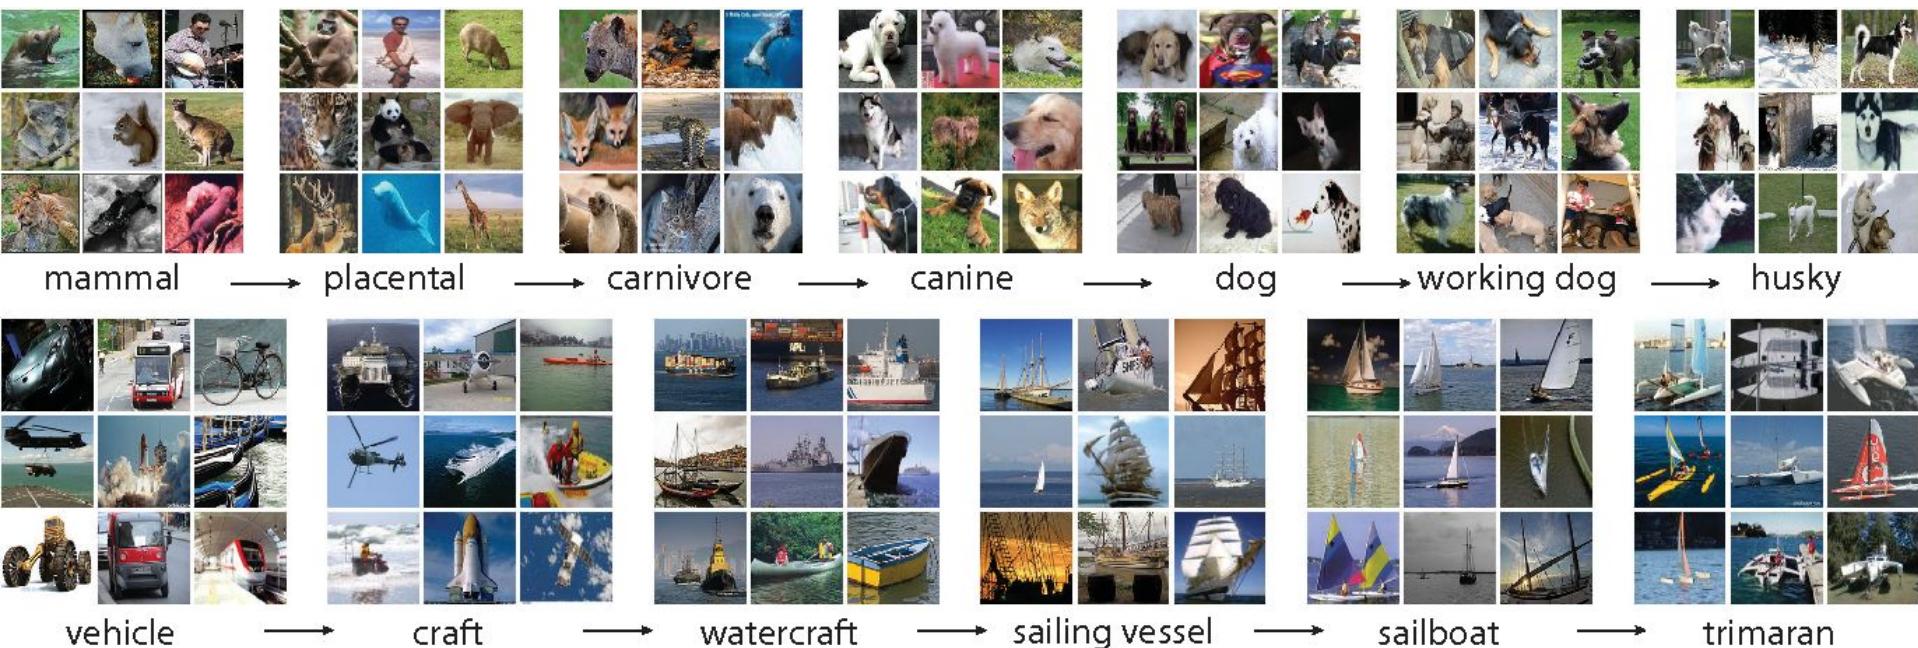
Second NN Winter

- NN cannot exploit many layers
 - Overfitting
 - Vanishing gradient (with NN training you need to multiply several small numbers → they become smaller and smaller)
- Lack of processing power (no GPUs)
- Lack of data (no large annotated datasets)
- Kernel Machines (e.g. SVMs) suddenly become very popular°



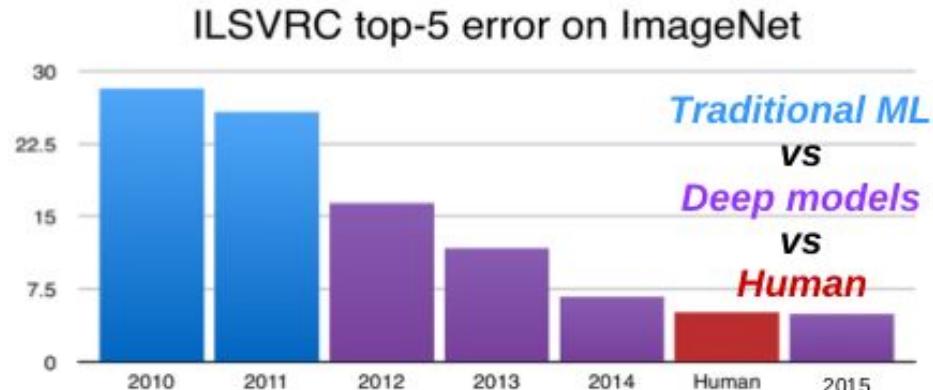
2009: ImageNet

A Large-Scale Hierarchical Image Database



2012: AlexNet

- Hinton's group implemented a CNN similar to LeNet [LeCun1998] but...
 - Trained on ImageNet (1.4M images, 1K categories)
 - With 2 GPUs
 - Other technical improvements (ReLU, dropout, data augmentation)



Convolutional NN (CNN)

- **Convolutional** layer: two functions produce a third that describes how the shape of one is changed by the other
- **pooling** layer: reduce dimensionality

Source layer

5	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	2	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Convolutional kernel

-1	0	1
2	1	2
1	-2	0

Destination layer

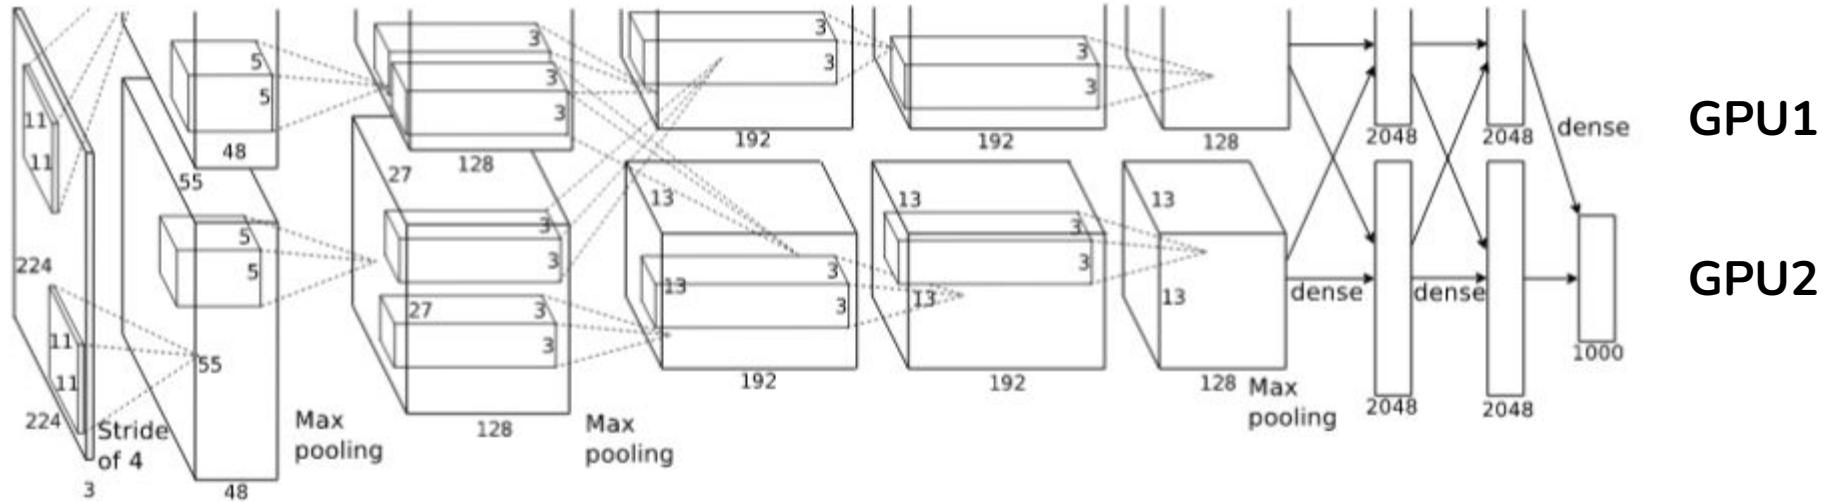
$$(-1 \times 5) + (0 \times 2) + (1 \times 6) + \\ (2 \times 4) + (1 \times 3) + (2 \times 4) + \\ (1 \times 3) + (-2 \times 9) + (0 \times 2) = 5$$

4	1	5	0
7	8	9	8
3	5	6	5
2	4	1	0

Max pooling

8	9
5	6

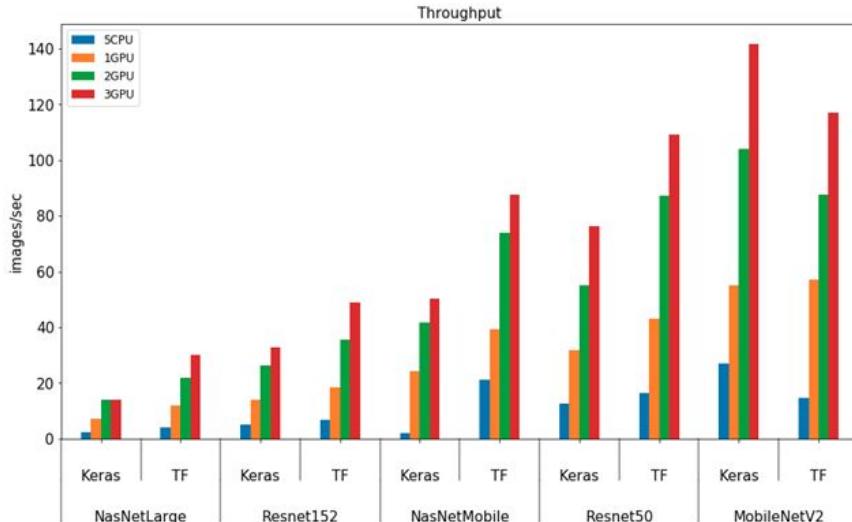
AlexNet



- 60M parameters
- Limited information exchange between GPUs

Why DL now?

- Three main factors:
 - Better hardware
 - Big data
- Technical advances:
 - Layer-wise pretraining
 - Optimization (e.g. Adam, batch normalization)
 - Regularization (e.g. dropout)



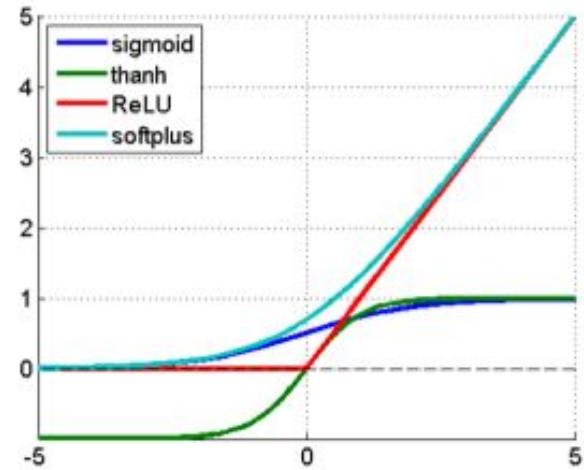
....

2010: Rectified Linear Units (ReLU)

$$f(x) = \max(0, x)$$

Activation function

- More efficient gradient propagation: (derivative is 0 or constant)
- More efficient computation: (only comparison, addition and multiplication).
- Sparse activation: e.g. in a randomly initialized networks, only about 50% of hidden units are activated (having a non-zero output)

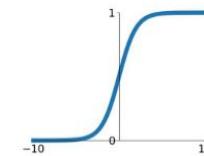


Activation functions

- Classification: sigmoid functions
 - sigmoids and tanh functions are sometimes avoided due to the vanishing gradient problem
- ReLU function is a general activation function
 - dead neurons in our networks -> the leaky ReLU
 - ReLU function should only be used in the hidden layers
 - As a rule of thumb, start with ReLU

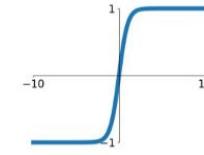
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



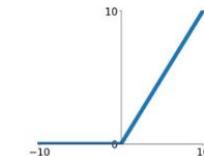
tanh

$$\tanh(x)$$



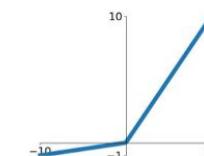
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

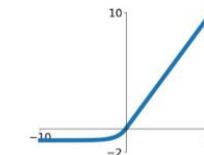


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Regularization

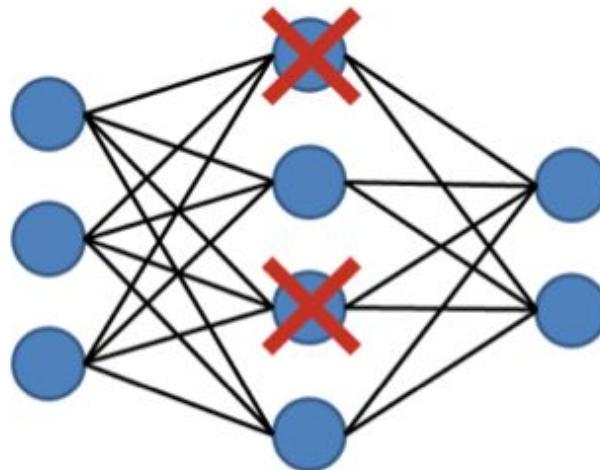
- One of the major aspects of training the model is overfitting -> the ML model captures the noise in your training dataset
- The **regularization** term is an addition to the loss function which helps generalize the model
 - **L1** or Lasso regularization adds a penalty which is the sum of the absolute values of the weights
 - L1+MSE
 - **L2** or Ridge regularization adds a penalty which is the sum of the squared values of weights
 - L2+MSE
- **Early Stopping** is a time regularization technique which stops training based on given criteria

$$\text{Min} \left(\sum_{i=1}^n (y_i - w_i x_i)^2 + p \sum_{i=1}^n |w_i| \right)$$

$$\text{Min} \left(\sum_{i=1}^n (y_i - w_i x_i)^2 + p \sum_{i=1}^n (w_i)^2 \right)$$

Dropout

- For each instance drop a node (hidden or input) and its connections with probability p and train
- Final net just has all averaged weights (actually scaled by $1-p$)
- As if ensembling 2^n different network substructures



Data augmentation

- Techniques to significantly increase the diversity of data available for training models, without actually collecting new data
- Data augmentation techniques such as cropping, padding, and horizontal flipping are commonly used to train large neural networks



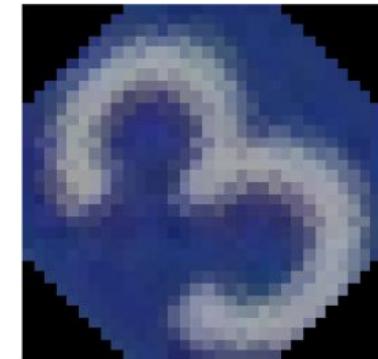
Original



Horizontal Flip

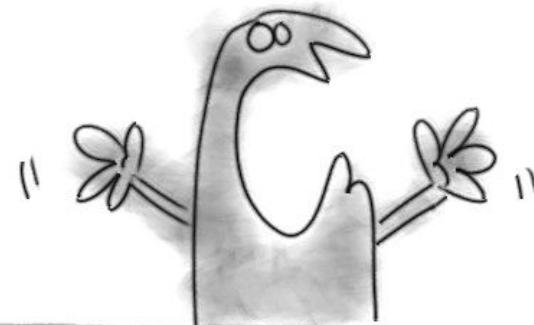


Pad & Crop



Rotate

Now What?!!



Training a neural network

- **Problem:** Predict gender from weight and height
- **Input Dataset:**

Name	Features		Labels
	Weight (lb)	Height (in)	
Alice	133	65	F
Bob	160	72	M
Charlie	152	70	M
Diana	120	60	F

1. Feature engineering

Name	Weight (lb)	Height (in)	Gender
Alice	133	65	F
Bob	160	72	M
Charlie	152	70	M
Diana	120	60	F

- Symmetrize numeric values
- Category -> numbers

Name	Weight (minus 135)	Height (minus 66)	Gender
Alice	-2	-1	1
Bob	25	6	0
Charlie	17	4	0
Diana	-15	-6	1



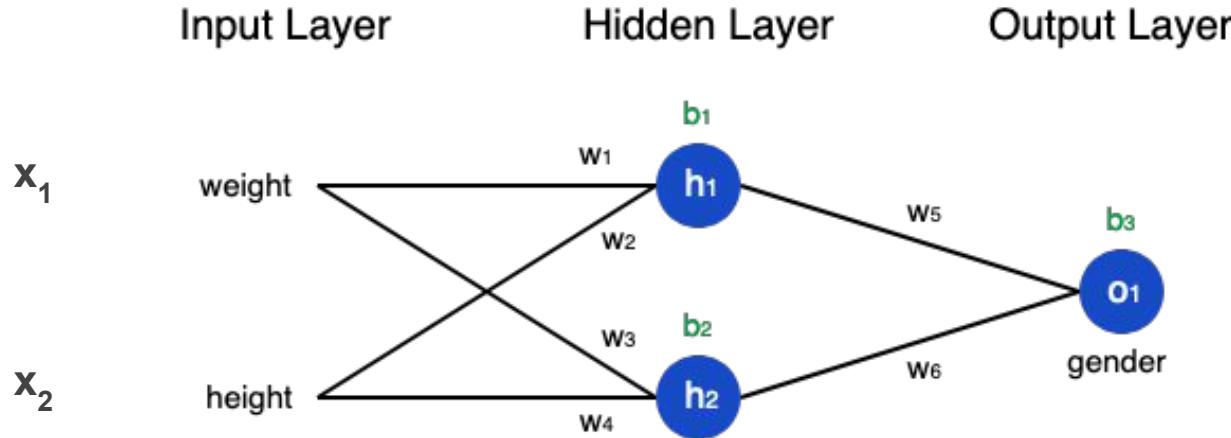
Ingredients

- n : 4, number of samples (Alice, Bob, Charlie, Diana)
- Inputs: X , dimension = 2
 - weights (X_1) and heights (X_2)
- Features: $x = (x_1, x_2)$, transformed inputs
- y : variable being predicted (Gender)
- y_{true} : true value of y

Name	Weight (minus 135)	Height (minus 66)	Gender
Alice	-2	-1	1
Bob	25	6	0
Charlie	17	4	0
Diana	-15	-6	1

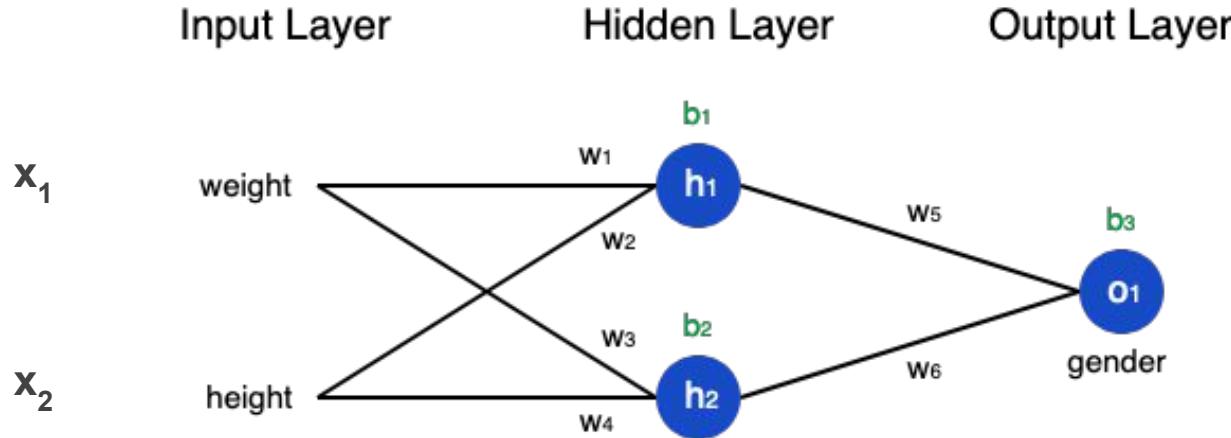
2. Model

- y_{pred} : predicted value of $y = \mathbf{o}$
- Neural network: with **1 hidden layer** with **2 neurons**
- Outputs of the hidden layer: \mathbf{h}
- Unknown parameters: weights \mathbf{w} and biases \mathbf{b}



Some math

- For each **neuron**: $y_j = b_j + \mathbf{f} \sum_i x_i w_{ij}$
 - **f** is the activation function
- For the **net**:
 - $h_1 = f(w_1 x_1 + w_2 x_2) + b_1$
 - $h_2 = f(w_3 x_1 + w_4 x_2) + b_2$
 - $o_1 = f(w_5 h_1 + w_6 h_2) + b_3$



3. Model training

- Training the network == Find weights \mathbf{w} and biases \mathbf{b} that minimize the loss
- Loss function \mathbf{L} : MSE
- Find weights \mathbf{w} and biases \mathbf{b} to minimise

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_{true} - y_{pred})^2$$

$$L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$$

- Typically, this is already implemented in ML software packages

1986: Back propagation

- Gradient descent
- Chain rule
- Layered Computation

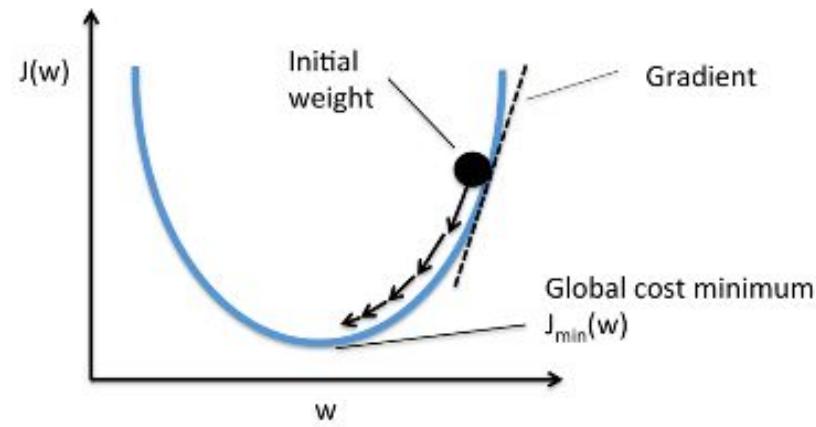
- Minimization taking **partial derivatives**
- For very simple case: with only Alice in the dataset, $n=1$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$$
$$MSE = \frac{1}{n} \sum_{i=1}^n (y_{true} - y_{pred})^2$$
$$L = (1 - y_{pred})^2$$
$$y_{pred} = o_1 = f(w_5 h_1 + w_6 h_2 + b_3)$$
$$h_1 = f(w_1 x_1 + w_2 x_2 + b_1)$$
$$\frac{\partial L}{\partial y_{pred}} = \frac{\partial(1 - y_{pred})^2}{\partial y_{pred}}$$
$$\frac{\partial y_{pred}}{\partial h_1} = w_5 * f'(w_5 h_1 + w_6 h_2 + b_3)$$
$$\frac{\partial h_1}{\partial w_1} = x_1 * f'(w_1 x_1 + w_2 x_2 + b_1)$$

Stochastic Gradient Descent (SGD)

- Start with randomly initialised weights
- update equation:

$$w_1 \leftarrow w_1 - \eta \frac{\partial L}{\partial w_1}$$



- η is a constant called the **learning rate** that controls how fast we train
- If $\frac{\partial L}{\partial w_1}$ is positive, w_1 will decrease, which makes L decrease.
- If $\frac{\partial L}{\partial w_1}$ is negative, w_1 will increase, which makes L decrease.

Stochastic vs Batch Gradient Descent (BGD) vs mini-batch

- In SVG the parameters are updated using only one single training instance (usually randomly selected) in each iteration vs the whole training set (== batch)
- Use mini-batch **sampled** in the dataset for gradient estimate.
- Sometimes helps to escape from local minima
- Noisy gradients act as regularization
- Variance of gradients increases when batch size decreases
- Not clear how many sample per batch

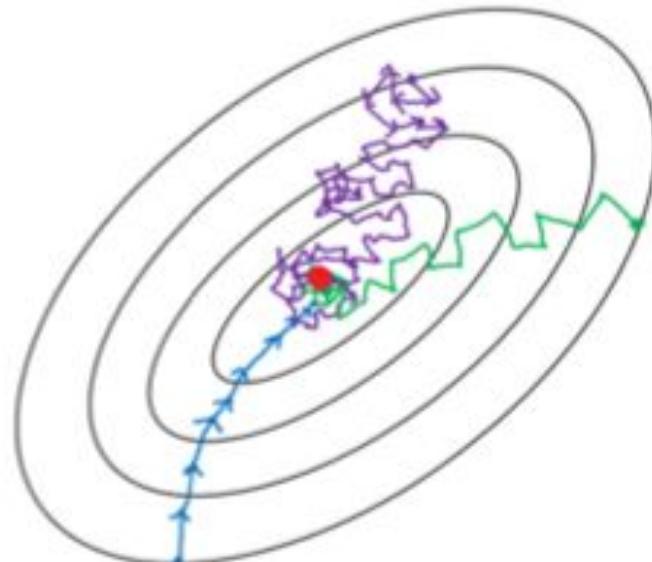
What happens in **one epoch** for SGD:

1. Take a random sample
2. Feed it to Neural Network
3. Calculate its gradient
4. Use the gradient we calculated in step 3 to update the weights
5. Repeat steps 1–4 for all the examples in training dataset

What happens in **one epoch** for Mini-BGD:

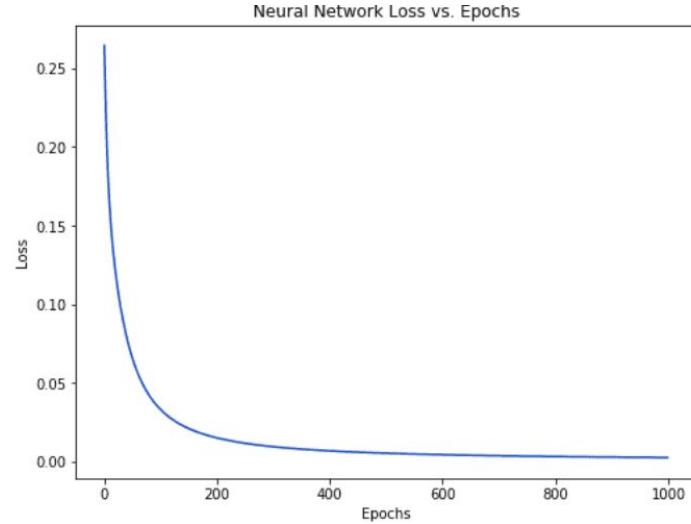
1. Pick a mini-batch
2. Feed it to Neural Network
3. Calculate the mean gradient of the mini-batch
4. Use the mean gradient we calculated in step 3 to update the weights
5. Repeat steps 1–4 for the mini-batches we created

- Batch gradient descent
- Mini-batch gradient Descent
- Stochastic gradient descent



Take home messages

- **Gradient descent** is an iterative learning algorithm that uses a training dataset to update a model.
 - **BGD**: Batch Size = Size of Training Set
 - **SGD**: Batch Size = 1
 - **Mini-BGD**. $1 < \text{Batch Size} < \text{Size of Training Set}$
- The **batch size** is a hyperparameter of gradient descent that controls the number of training samples to work through before the model's internal parameters are updated.



- The number of **epochs** is a hyperparameter of gradient descent that controls the number of complete passes through the training dataset