

Big data science

Day 5

F. Legger - INFN Torino

<https://github.com/Course-bigDataAndML/MLCourse-2425>

- So long, and thanks for all the images!
 - Taken freely from the web
 - Credits go to original creators



We Learned

- Big data
 - Analytics
 - Machine learning
 - Deep learning

Today



- NN models
 - Parallelisation
 - Heterogeneous architectures

Recap:

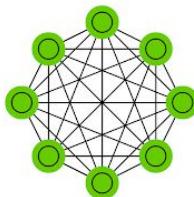
- **Machine learning (ML)**: family of algorithms with ability to automatically learn and improve from experience without being explicitly programmed
 - potential to approximate linear and non-linear relationships
- For a given problem:
 - *Choose architecture*
 - *Train model*
 - *Tune hyperparameters*
 - *Do cross-validation*
 - *Do inference*

Neural Networks

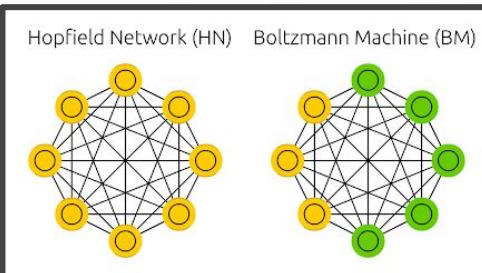
©2019 Fjodor van Veen & Stefan Leijnen asimovinstitute.org

- Input Cell
- Backfed Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Capsule Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Gated Memory Cell
- Kernel
- Convolution or Pool

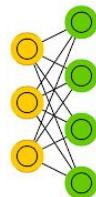
Markov Chain (MC)



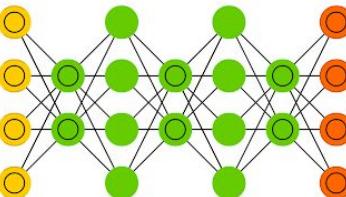
Hopfield Network (HN) Boltzmann Machine (BM)



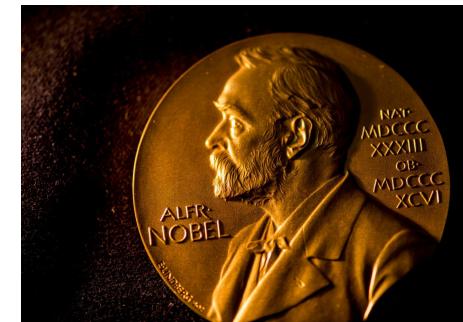
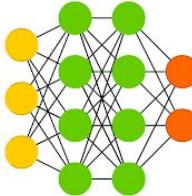
Restricted BM (RBM)



Deep Belief Network (DBN)

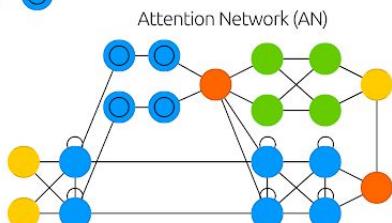
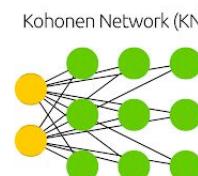
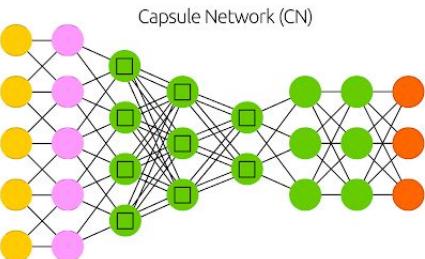
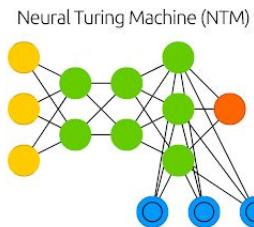
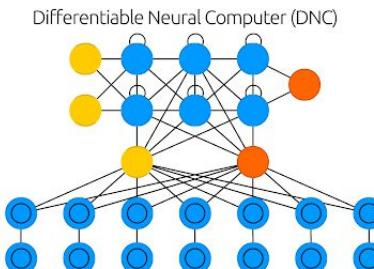
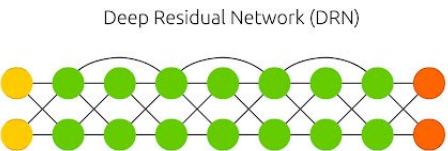
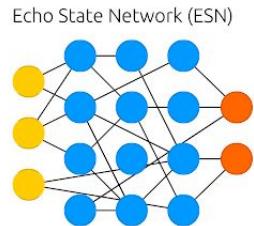
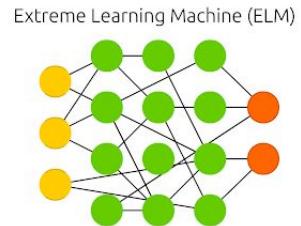
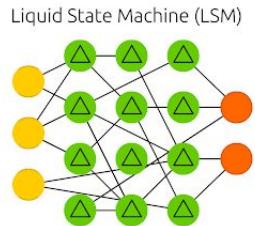
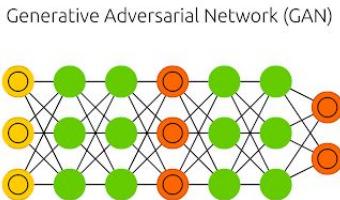
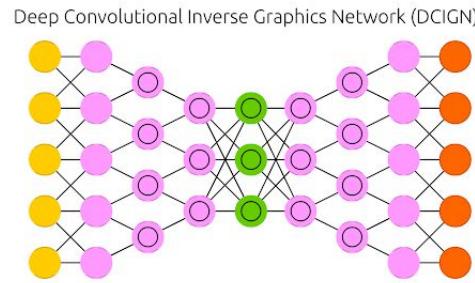
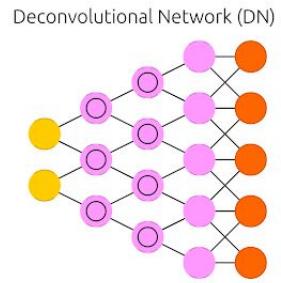
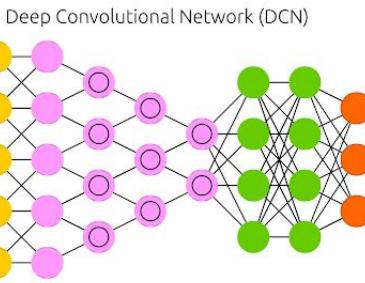


Deep Feed Forward (DFF)



Nobel Prize
Physics 2024

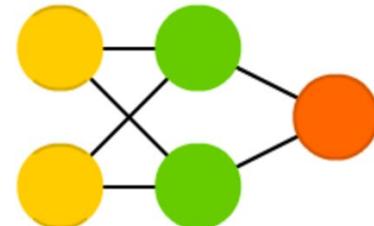
- Input Cell
- Backfed Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Capsule Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Gated Memory Cell
- Kernel
- Convolution or Pool



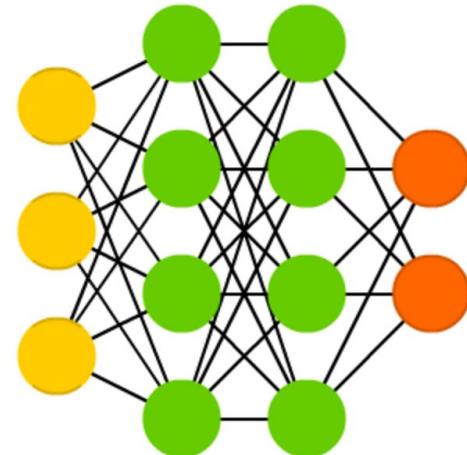
Feed Forward

- Supervised, simplest form of NN
 - Used in many ML tasks, speech, image recognition, classification, computer vision
 - Easy to implement and combine with other type of ML algorithms
- input/outputs are vectors of fixed length
- data passes through input nodes and exits on the output nodes
- typically trained with back-propagation
- **DFF** is a FF NN with more than one hidden layer

Feed Forward (FF)

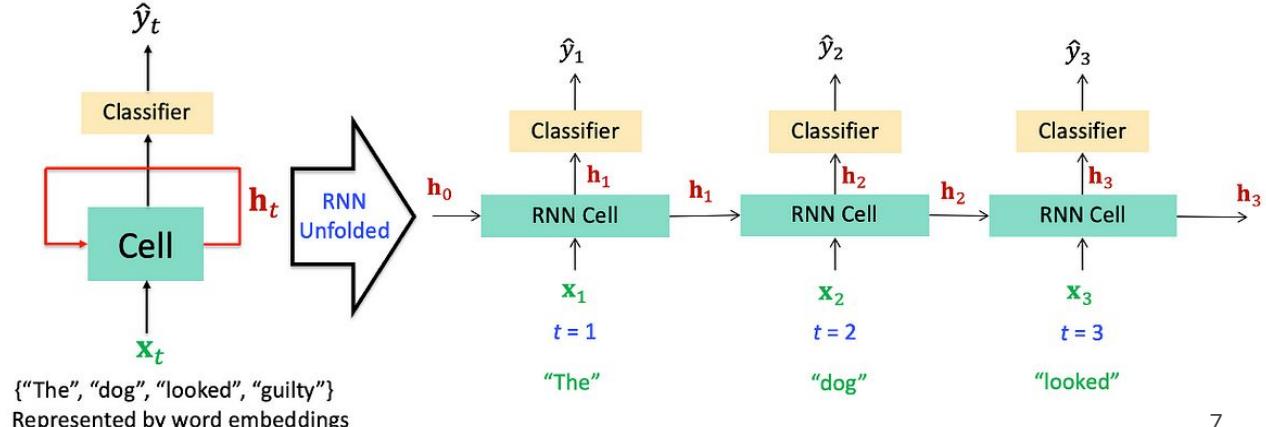


Deep Feed Forward (DFF)



Recurrent Neural Network (RNN)

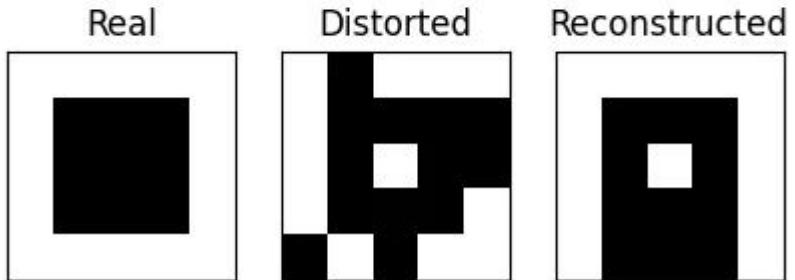
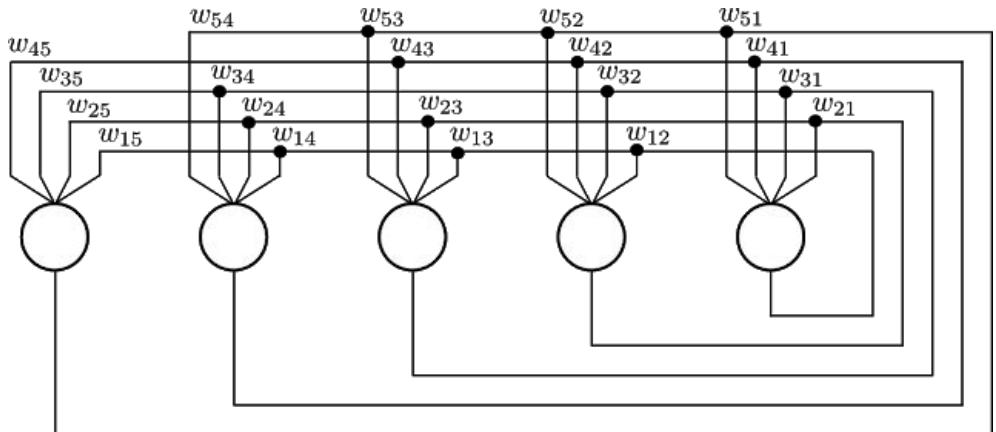
- FFNN with **Recurrent Cells**: cell that receives its own output with fixed delay
- **RNNs** permit to operate on **sequences of vectors**
 - context is important, decision from past iterations can influence current state
 - a word can be analyzed only in context of previous words or sentences
- RNNs, once unfolded in time, can be seen as **very deep FF networks** in which all the layers share the same weights
 - The parameters to be learned are shared by all time steps



Hopfield Network

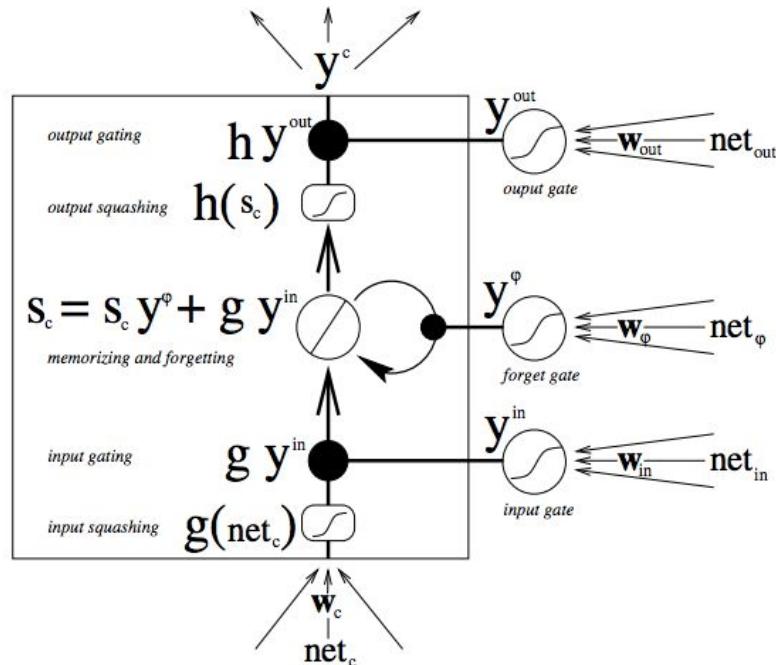
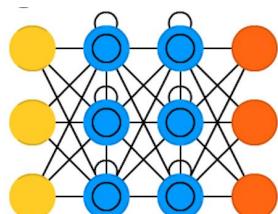


- A Hopfield network (or **associative memory**) is a form of RNN
- It consists of a single layer of neurons, where each neuron is connected to every other neuron except itself
- the connections are bidirectional and symmetric



Long/Short Term Memory (LSTM)

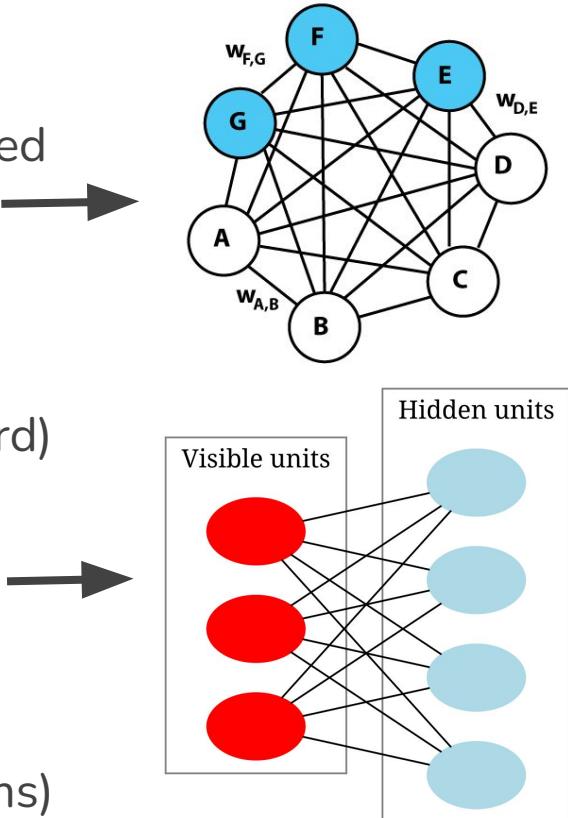
- RNNs not really capable of learning long term dependencies
 - Due to vanishing gradient with increasing time steps
- A common LSTM unit is made of a **cell**, an **input**, an **output** and a **forget** gate
 - The cell remembers values over arbitrary time intervals and the three gates regulate the flow of information into and out of the cell



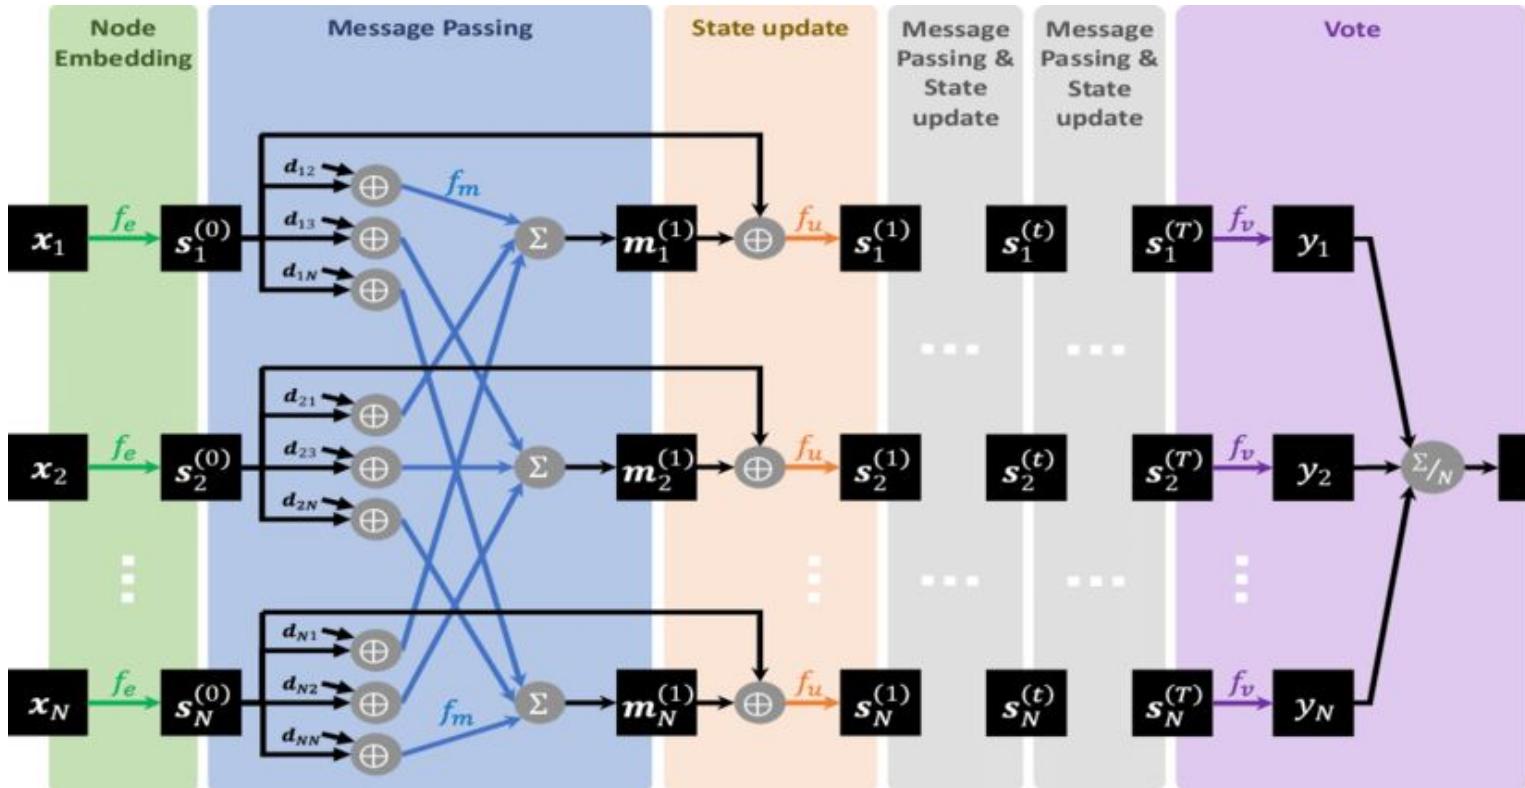
Boltzmann machine



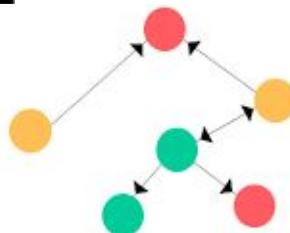
- Neurons from input layer and hidden layers are connected
 - Learns only trivial relationships
- **Restricted Boltzmann machines (RBM):**
 - No connections between units of the same type
 - connections going both ways (forward and backward) that have a probabilistic / energy interpretation
- RBMs are primarily used for:
 - Feature extraction
 - Dimensionality reduction
 - Collaborative filtering (e.g., recommendation systems)



Graph Neural Network (GNN)

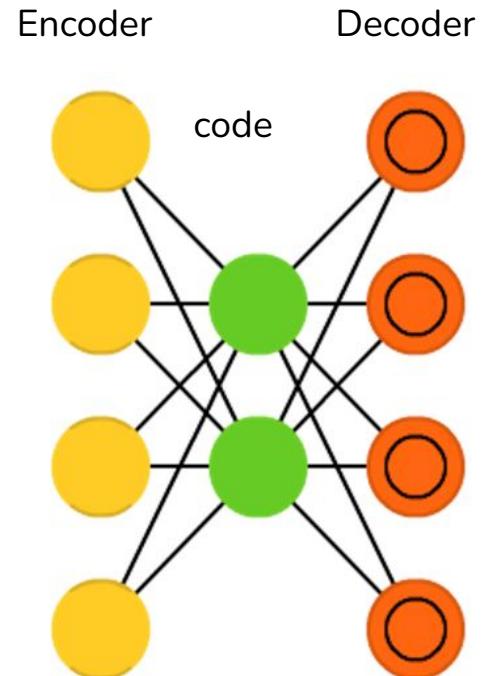


- Node prediction
- Edge prediction
- Graph prediction



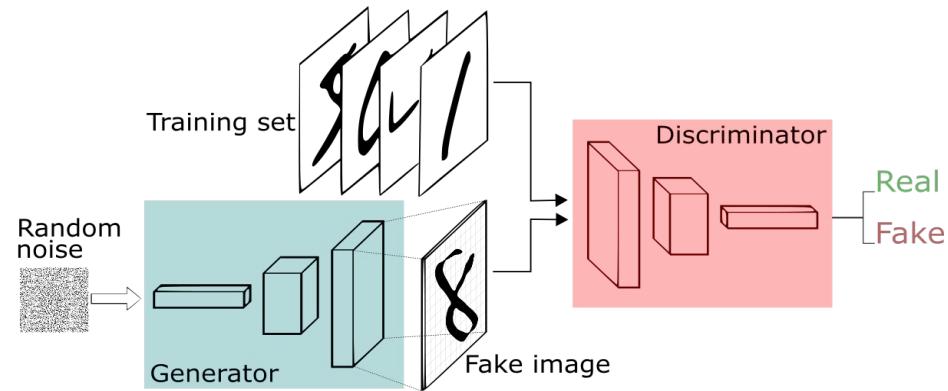
Auto Encoders

- used for classification, clustering and feature compression (**unsupervised learning**)
- Compress (encode) information automatically
 - An **encoder** is a deterministic mapping f that transforms an input vector x into hidden representation y
 - A **decoder** maps back the hidden representation y to the reconstructed input z via g
- **Autoencoder:** compare the reconstructed input z to the original input x and tries to minimize the reconstruction error



Generative Adversarial Networks (GAN)

- GANs represent a huge family of double networks that are composed from a **generator** net and a **discriminator** net
 - The generator produces samples close to training samples
 - Discriminator net (adversary) differentiates between samples from the generative net and the training set
 - Use error feedback to improve task of both nets, until discriminator can no longer distinguish
- Can be used to generate samples of data without prior knowledge of the data

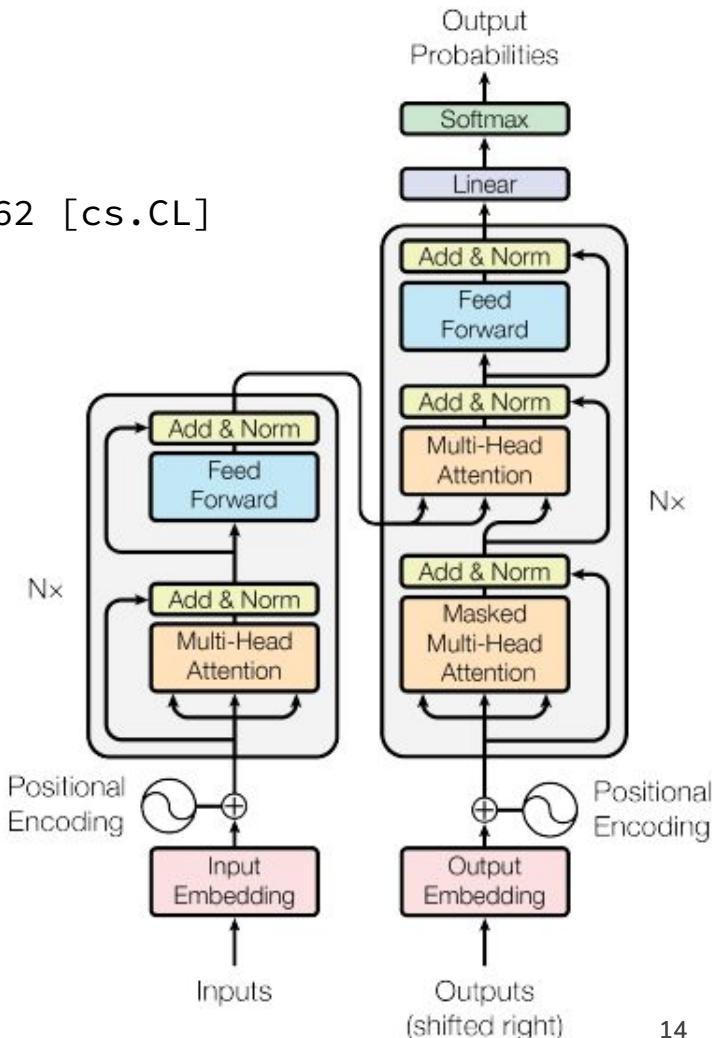


Transformers

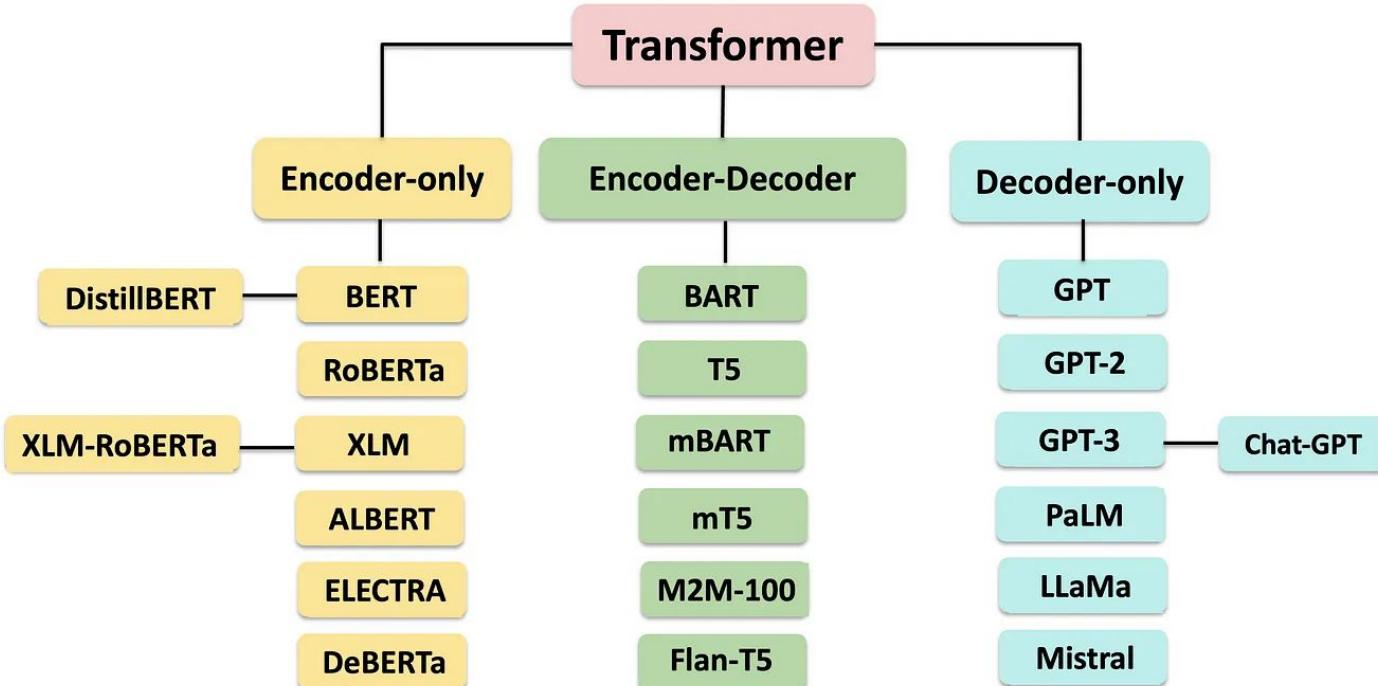
arXiv:1706.03762 [cs.CL]

- All you need is **attention**
 - transformers process the entire input all at once
 - the attention mechanism provides context for any position in the input sequence
- Self-attention: **query, key, value:**
 - The significance of each part of the input data is differentially weighted

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

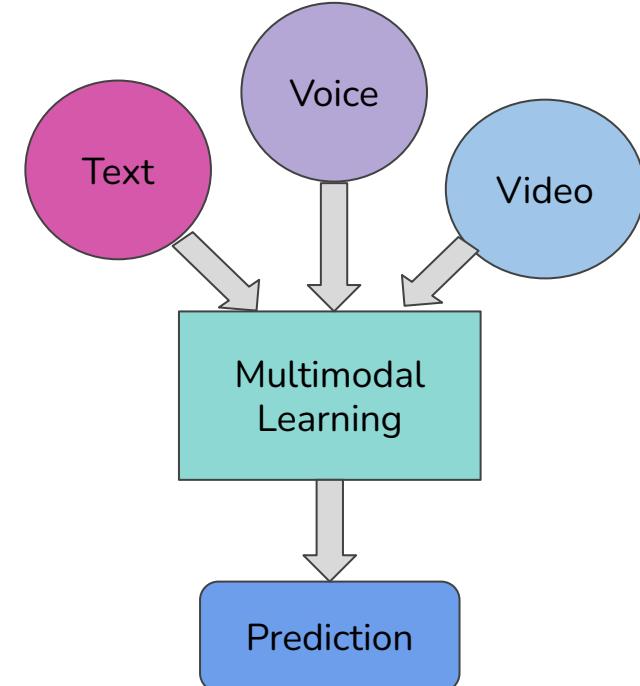


A world of Transformers



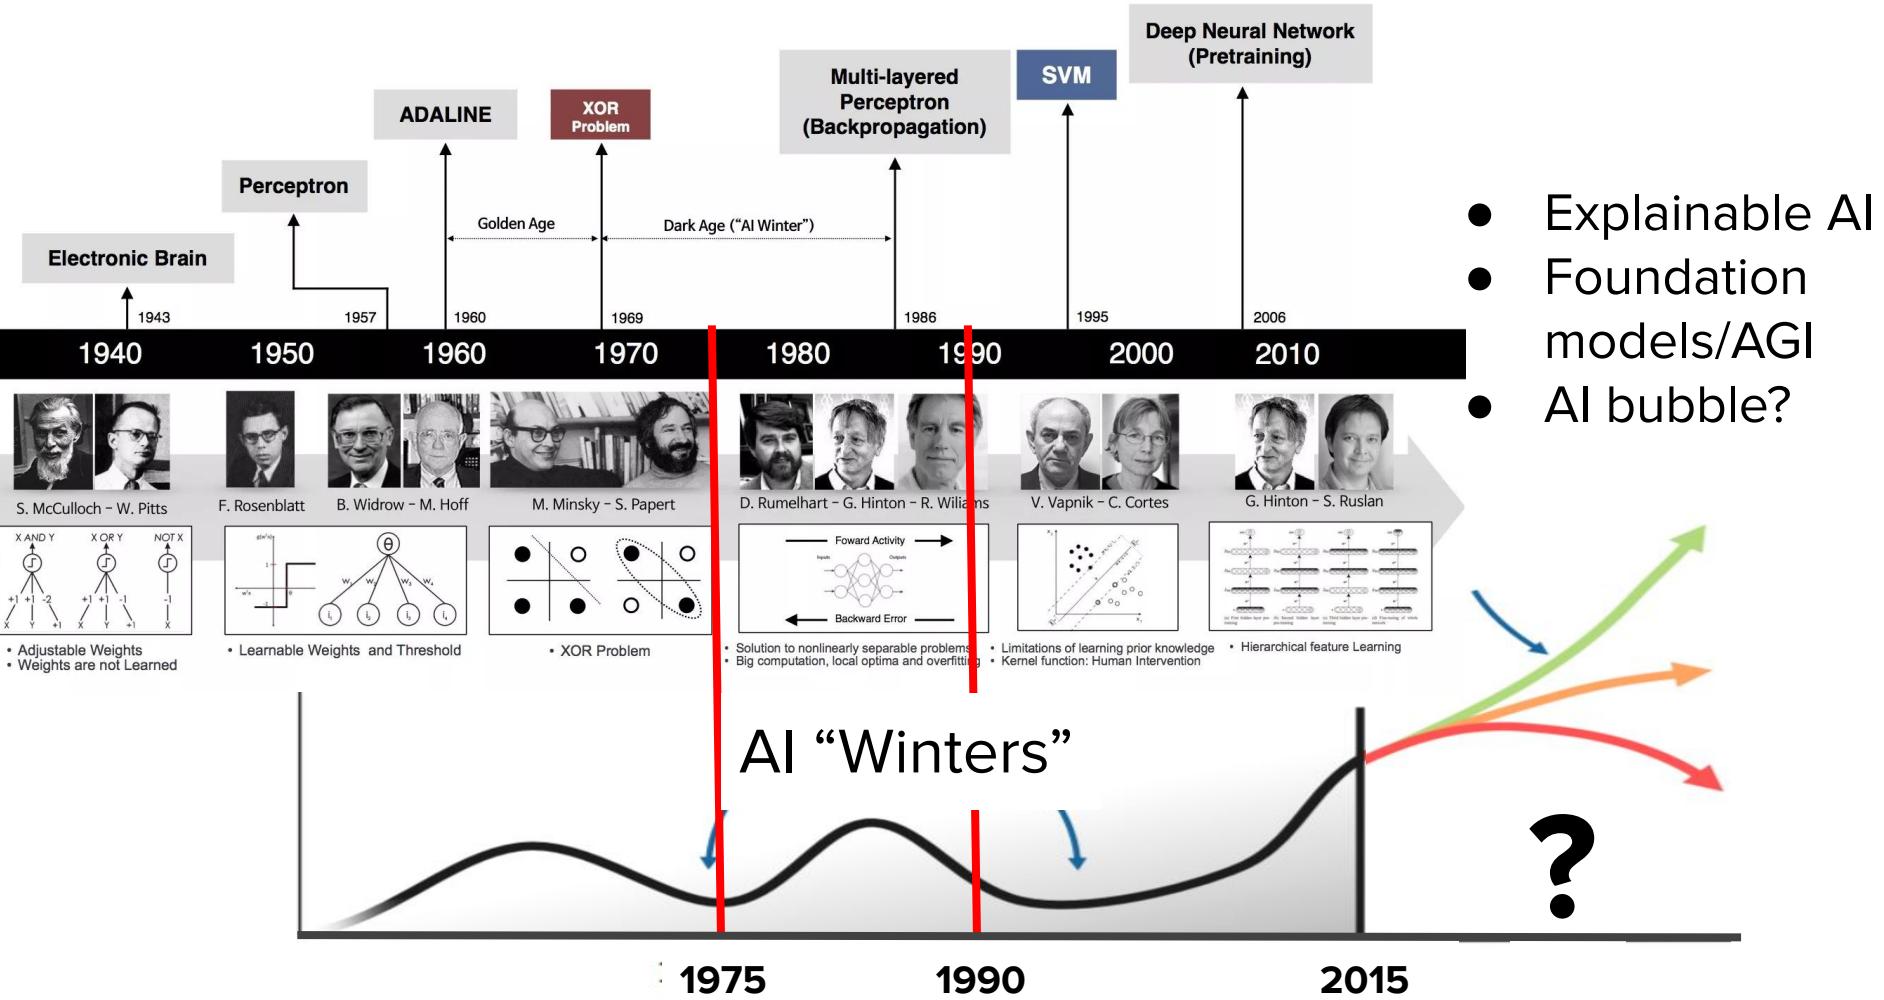
Multimodal Learning

- A **multimodal foundation model (MFM)** is a type of generative AI that can process and output multiple data types
- Learning techniques:
 - **Fusion:** encode the different modalities into a common representation space
 - **Alignment:** create modality-invariant representations that can be compared across modalities
 - **Late Fusion:** combine predictions from models trained on each modality separately



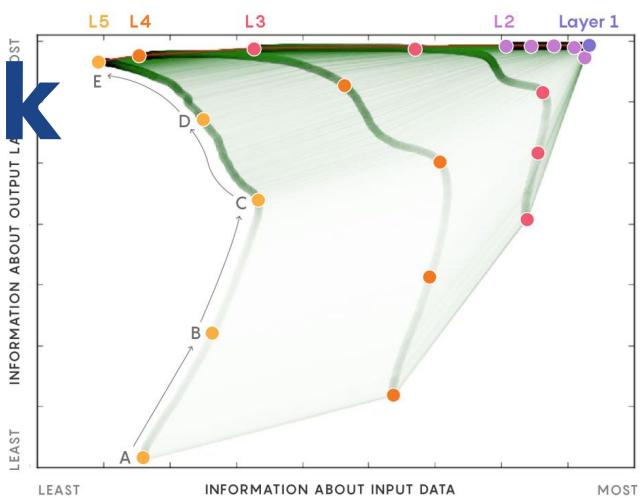
NN evolution in past 15 years

- 2009: handwriting recognition prizes with LSTM
- 2011 DanNet (CNN) beats humans in **visual pattern recognition**
- 2014: GANs, Alexa, Tesla AutoPilot
- 2015: FaceNet starts facial recognition programs
- 2016: AlphaGo **beats Go champion**
- 2017: Google **Translate** and Facebook Translate based on two LSTMs
- 2018: Cambridge analytica, deep fakes
- 2019: **DeepMind** defeats professional StarCraft players with RL + LSTM, Rubik's Cube solved with a Robot Hand
- 2021: AI Colorectal Cancer Detection Technique Better than Pathologist
- 2022: LLMs such as **chatGPT** produce “believable” open text
- 2023: LLMs explosion: **LLaMA, BARD**
- 2024: First MMLLM: GPT-4, Gemini
- 2025: AI generated song tops **Spotify’s Viral 50** songs in the US



Information Bottleneck

- based on the notion of minimal sufficient statistics for extracting information about the target
- trade-off between the complexity of representation and the power of predicting
- In the first epochs, the network is trained to fully represent the input data; then, it learns to forget the irrelevant details by compressing the representation of the input



A INITIAL STATE: Neurons in Layer 1 encode everything about the input data, including all information about its label. Neurons in the highest layers are in a nearly random state bearing little to no relationship to the data or its label.

B FITTING PHASE: As deep learning begins, neurons in higher layers gain information about the input and get better at fitting labels to it.

C PHASE CHANGE: The layers suddenly shift gears and start to "forget" information about the input.

D COMPRESSION PHASE: Higher layers compress their representation of the input data, keeping what is most relevant to the output label. They get better at predicting the label.

E FINAL STATE: The last layer achieves an optimal balance of accuracy and compression, retaining only what is needed to predict the label.

Self Supervised Learning

- Supervised learning needs many labeled data
- Reinforced learning:
 - Not practical to train in real world (when no simulation is available)
 - takes longer than an average human for a machine to learn a new task
- **Self supervised learning:** Predict everything from everything else - learn representations, rather than learning specific tasks
 - Very large networks trained with large amount of data
 - Filling the blanks - Word2Vec, Transformer architecture for NLP
 - Not (yet) so successful for continuous problems (image, video)

Consciousness Prior

- Current deep learning:
 - **System 1:** fast, unconscious task solving
- Future deep learning:
 - **System 2:** slow, conscious task solving like reasoning, planning
- How?
 - Learn by predicting in abstract space
 - Learn representations (low dimensional vector), derived using attention from a high dimensional vector
 - The prior: the factor graph (joint distribution between a set of variables) is sparse

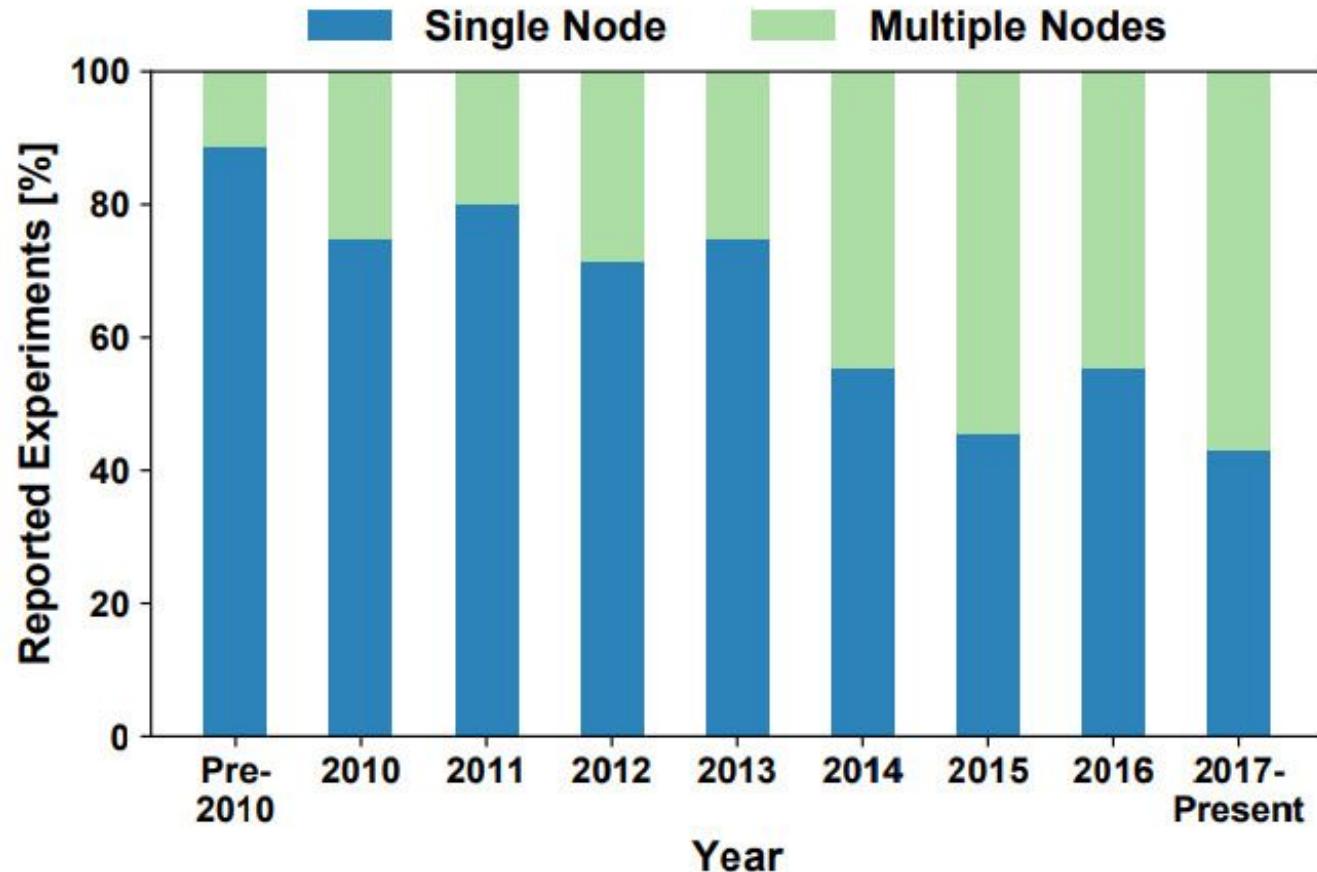
$$P(S) = \frac{\prod_j f_j(S_j)}{Z}$$

Artificial General Intelligence

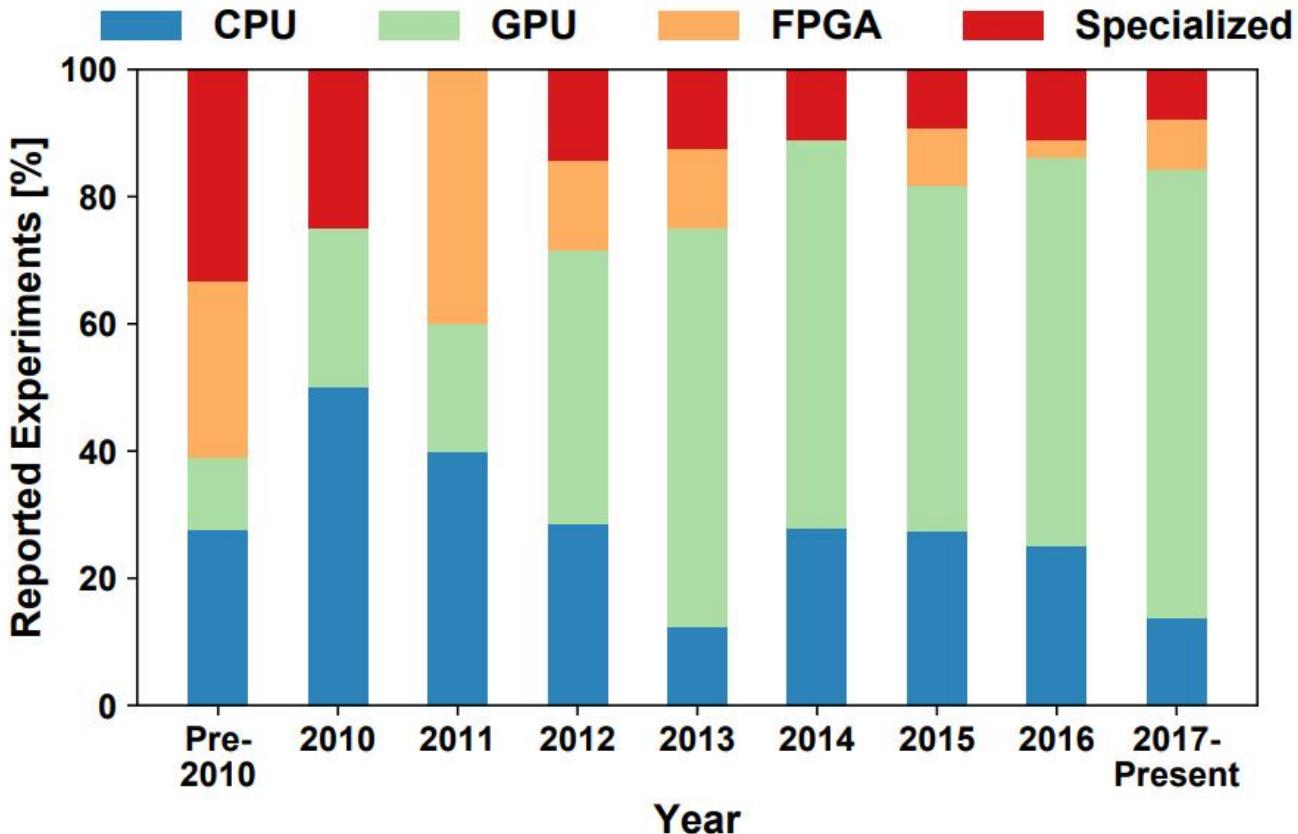
- **Common sense:**
 - Current systems may be easily fooled by just slight changes in the input data (for example image taken from another viewpoint)
 - Embed coordinate systems, whole-part relationship (capsules)
- **Abstract concepts:**
 - Current models may be able to distinguish between a jet and a tau, but do not know what a particle is
- **Generalisation/Creativity:**
 - Current models highly specialised and engineered to solve specific problems

- Most ML algorithms require significant amount of CPU, RAM and sometimes GPU in order to be applied efficiently
- **Does it fit on your laptop?**





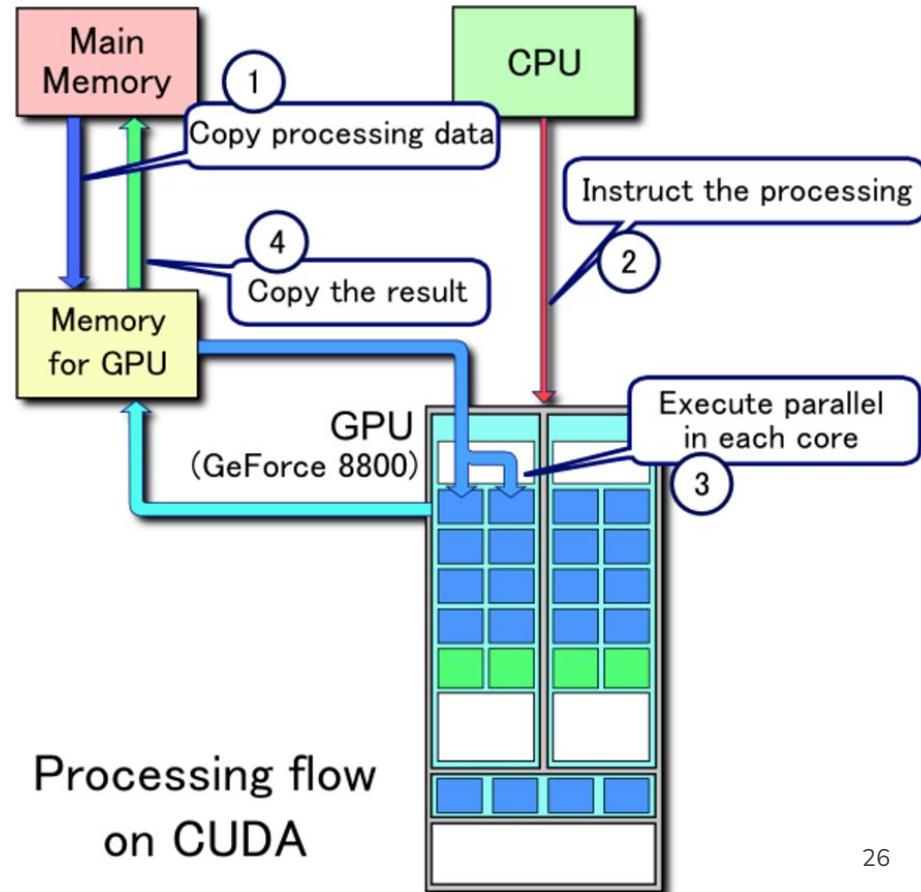
Fraction of non-distributed vs distributed deep learning over time



Use of hardware for deep learning

GPU (Graphical Processing Unit)

- Typically composed of thousands of logical cores
- Excels at matrix and vector operations
 - Gaming, rendering...and ML
- Main vendor: NVidia
 - CUDA: parallel computing platform and programming model



Options

- **NVIDIA CUDA**
 - Market leader, large user community, best support from DL frameworks
 - Can be used in python, C/C++, fortran, Matlab
- **AMD HIP**
 - Limited support for pyTorch and Tensorflow
 - Slightly behind in terms of performances
- **INTEL**
 - Intel Arc
 - (Kind of) support for Tensorflow/pytorch
- **Google TPU**
 - Good performances, more powerful than cloud GPUs
 - best for training, for prototype and inference better use cheaper alternatives
- **Amazon AWS and Microsoft Azure**
 - Powerful, easy to scale, expensive

How to choose hardware

- GPUs (may) have high performance for floating point arithmetic but
 - limited amount of memory available
 - rate at which data can be moved from CPU to GPU
 - **memory bandwidth for GPUs must be seen relative to the amount of FLOPS:** if one has more floating point units, a higher bandwidth is needed to keep them occupied
- Be aware of power consumption and overheating when placing multiple GPUs close to one another!

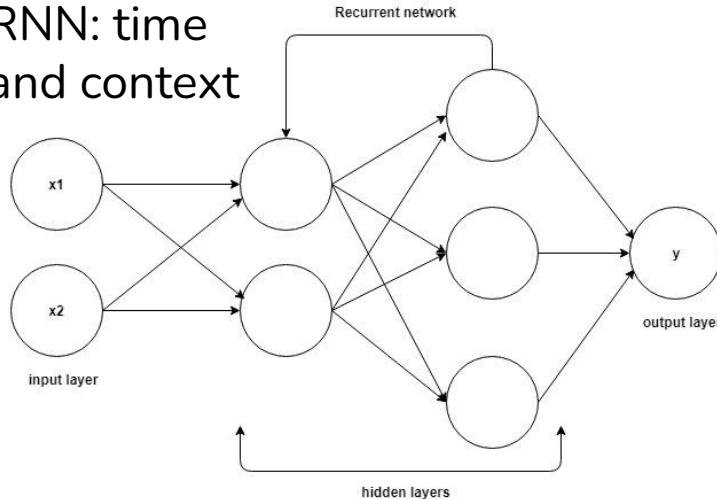
How to choose hardware for NN

- CNNs need a lot of computing power (need high number of cores, FLOPS)
 - Remember AlexNet?
- Matrix operations: typically bandwidth cost is larger than multiplication cost
 - Particularly important for many small multiplications, such as those required to train **LSTMs** and **RNNs**

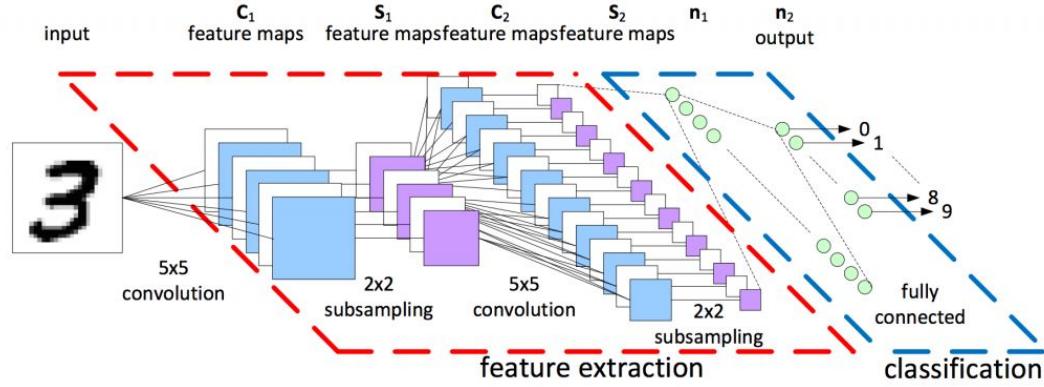
Parallelisation on multiple GPUs

- ‘Easy’ for RNNs and CNNs
- Fully connected networks with transformers poorer performances
- multiple GPUs can be used for tasks trivial to parallelise such as hyperparameter scan

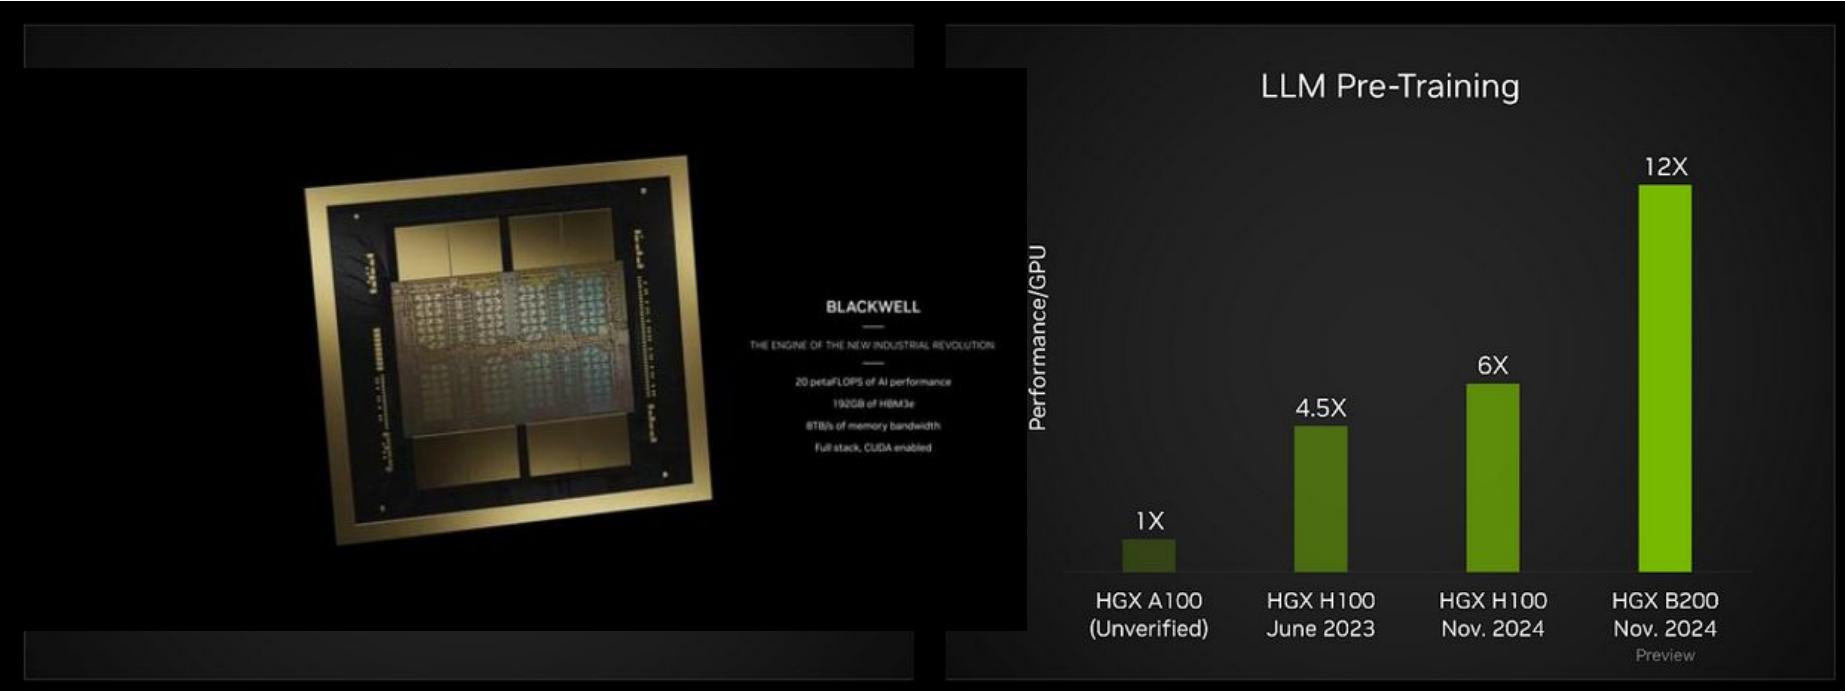
RNN: time
and context



CNN: image recognition



GPU partitioning



With its multi-instance GPU (MIG) technology, A100 can be partitioned into up to seven GPU instances, each with 10GB of memory.

Local training

- Both model and data fit on a **single machine (multi cores + GPU)**
 - Multi-core processing:
 - *Embarrassingly parallel process*: use the cores to process multiple images at once, in each layer
 - Use multiple cores to perform SGD of multiple mini-batches in parallel.
 - Use GPU for computationally intensive subroutines like matrix multiplication.
 - Use both multi-core processing and GPU where all cores share the GPU and computationally intensive subroutines are pushed to the GPU.

Distributed training

- Data or model stored across **multiple machines**
- **Data parallelism:** data is distributed across multiple machines
 - data is too large to be stored on a single machine or to achieve faster training
 - large batch of input data split across a collection of workers, each holding the full model
 - the forward pass involves no communication
 - the backward pass involves aggregating the gradients computed by each individual worker with respect to its separate part of the “global batch”

Data parallelism

- **Synchronous update:** all loss gradients in a given mini-batch are computed using the same weights and full information of the average loss in a given mini-batch is used to update weights
- **Asynchronous update:** as soon as a machine finishes computing updates, the parameters in the driver get updated. Any machine using the parameters will fetch the updated parameters from the server.
- **scaling out:** the “global batch size” (i.e. the total number of samples across all workers that are seen in a single forward pass) increases
 - Affects convergence of DL algorithms and does not reach same accuracy levels that can be obtained with smaller batch sizes

Model parallelism

- model layers split across a collection of workers
- the batch size stays constant, and large models that would not fit the memory of a single device can be trained
- active communication also in the forward pass, thus requiring a lot more communication than a data parallel approach
- unless the advantages of model parallelism are absolutely critical (model doesn't fit on one machine), most research on large-scale training is done using data parallelism.
- Schemes involving a mix of data and model parallelism also exist: *hybrid parallelism*.

Other issues - not only compute

- **Communication** bottlenecks during gradient-aggregation, due to high ratio between the number of parameters and amount of computation
- **Memory** bottlenecks for large networks, particularly when using memory-limited GPUs
- **I/O contention** with large data sets, or data sets with many small samples, and the data has to be physically stored on shared storage facilities
 - On large system such as **HPCs**, data are typically stored on shared file systems
 - Might be faster to copy data locally to worker node

Architectural choices

- Your model doesn't fit in the GPU memory?
 - tune the model e.g. by reducing its connectivity (if that is acceptable) to fit the hardware
 - use model parallelism to distribute the model over multiple GPUs
 - use a data parallel approach on CPUs (if it does fit in CPU memory)
- Your dataset doesn't fit in the GPU memory?
 - Make sure your data arrives fast enough to the GPU, otherwise revert to CPU
- Do you have access to a system with a few GPU nodes, but thousands of CPUs?
 - Development stage: run a limited number of examples (potentially even with a smaller model) on a single GPU, which allows quick development cycles
 - Production: may require CPUs because of memory constraints

That's all Folks!

Spare slides



Requirements for ML framework

- Run any ML algorithm of our choice in parallel
- Large datasets to be processed
- Multi-tenancy, so different users can request simultaneously ML pipelines
- An efficient resource management system for the cluster
- Support heterogeneous architectures
 - CPUs, GPUs, ...

ML as a Service (MLaaS)

CHALLENGES

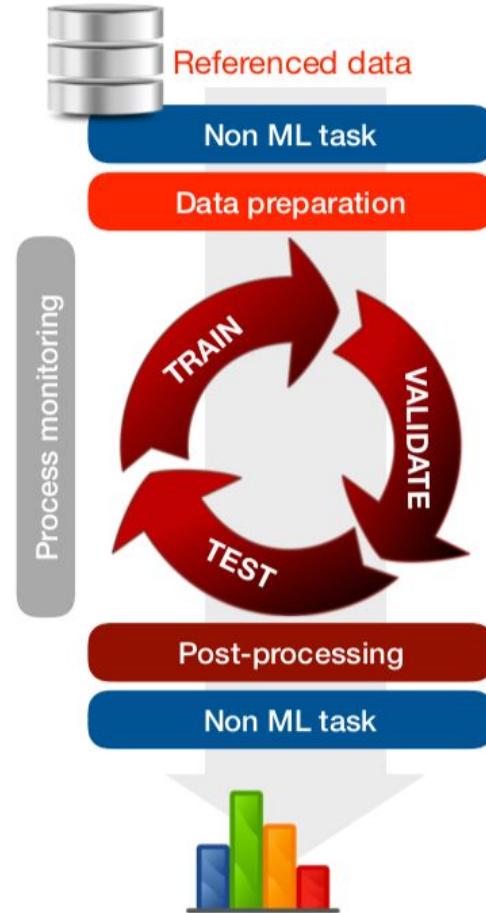
- Reconstruction
- Analysis
- Trigger
- Data quality
- Detector monitoring
- Computing operations
- Monte Carlo tuning
- ...

REQUIREMENTS

- Workflow definition
 - Results reproducibility
- Multi-tenancy (scheduling, authentication...)
- Parallel execution and scaling
- Data handling
- Ease of use and management
- ...

IMPLEMENTATION

- Lightweight virtualization
- Modularity
- Flexibility
- Heterogeneous back-end infrastructures
- ...



Example architecture

